

CS3281 / CS5281

# Process Creation and Control

CS3281 / CS5281

Spring 2026

*\*Some lecture slides borrowed and adapted from CMU's  
"Computer Systems: A Programmer's Perspective"*

# Review

- System calls are how user-level processes request services from the kernel
  - Many kinds of system calls: connect to a network host, read bytes, open a file, close a file, etc
- System calls are supported by special machine-code instructions
  - On x86: int 0x80 or syscall
  - On RISC-V: ecall
  - These “trapping” instructions cause a lot of hidden work to happen
    - Control flow jumps to the OS kernel
    - The kernel handles the request
    - Control returns to the user-space application
- Today we'll look at the system calls for:
  - Creating a process
  - Running a program
  - Waiting for a process to terminate

# System Call Error Handling

- On error, Libc system-level functions typically return -1 and set global variable errno to indicate cause
- Hard and fast rule:
  - You must check the return status of every system-level function
  - Only exception is the handful of functions that return void. E.g., exit().
- Example:

```
if ((pid = fork()) < 0) {
    fprintf(stderr, "fork error: %s\n", strerror(errno));
    exit(1);
}
```

# Obtaining Process IDs

## Linux

- `pid_t getpid(void)`
  - Returns PID (Process ID) of current process
- `pid_t getppid(void)`
  - Returns PID of parent process

Note that `pid_t` is just a signed 32 bit integer on most platforms.  
Type is defined for portability.

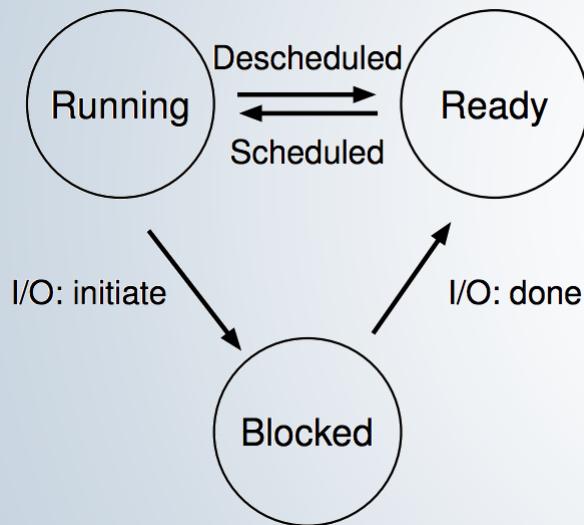
# Creating and Terminating Processes

From a programmer's perspective, we can think of a process as being in one of three states

- Running
  - Process is either executing on CPU
- Ready
  - Process waits to be executed and will eventually be *scheduled* (i.e., chosen to execute) by the kernel
  - i.e. process enters the ready state because of the context switch
- Stopped (blocked)
  - Process execution is *suspended* and will not be scheduled until further notice
  - i.e. process is blocked because it requests a resource that is not available now
- Terminated
  - Process is stopped permanently

# Processes don't run all the time

The three basic process states:



- OS *schedules* processes
  - Decides which of many competing processes to run.
- A *blocked* process is not ready to run.
- I/O means input/output – anything other than computing.
  - For example, reading/writing disk, sending network packet, waiting for keystroke, updating display.
  - While waiting for results, the process often cannot do anything, so it **blocks**, and the OS schedules a different process to run.

# Process State – Linux

- The kernel has a *process descriptor* (also called process control block or PCB) of type struct task\_struct for each process
  - Defined in <linux/sched.h>
- Process descriptor contains all the information about a process
- The kernel stores the list of processes in a circular doubly linked list called the task list
- What does the state of a process include?
  - State: running, ready, terminated, waiting
  - Priority
  - Parent
  - PID (process identifier)
  - Address space
  - Pending signals
  - Open files
  - Saved registers
  - etc.

# Terminating Processes

- Process becomes terminated for one of three reasons:
  - Receiving a signal whose default action is to terminate (next lecture)
  - Returning from the `main` routine
  - Calling the `exit` function
- `void exit(int status)`
  - Terminates with an *exit status* of `status`
  - Convention: normal return status is 0, nonzero on error
  - Another way to explicitly set the exit status is to return an integer value from the `main` routine
- `exit` is called **once** but **never** returns.

# Creating Processes: fork()

- *Parent process* creates a new running *child process* by calling `fork`
- `int fork(void)`
  - Returns 0 to the child process, child's PID to parent process
  - Child is *almost* identical to parent:
    - Child gets an identical (but separate) copy of the parent's virtual address space
    - Child gets identical copies of the parent's open file descriptors
    - Child has a different PID than the parent
- `fork` is interesting (and often confusing) because it is called *once* but returns *twice*

# Fork Example

- Call once, return twice

```
int main()
{
    int pid;
    int x = 1;

    pid = fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}
```

# Fork Example

- Call once, return twice
- Concurrent execution
  - Can't predict execution order of parent and child
- Duplicate but separate address space
  - x has a value of 1 when fork returns in parent and child
  - Subsequent changes to x are independent
- Shared open files

```
int main()
{
    int pid;
    int x = 1;

    pid = fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}
```

```
linux> ./fork
parent: x=0
child : x=2
```

# Question

- What does the following code do?

```
#include <stdio.h>
#include <sys/types.h>

int main() {
    while(1) {
        fork();
    }
    return 0;
}
```

# Question

- What does the following code do?

```
#include <stdio.h>
#include <sys/types.h>

int main() {
    while(1) {
        fork();
    }
    return 0;
}
```

- Creates a new process
  - Then each process creates a new process
  - Then each of those creates a new process...
- Known as a Fork bomb!
  - Machine eventually runs out of memory and processing power and will stop working
- Defense: limit number of processes per user

# Modeling fork() with Process Graphs

- A process graph is a useful tool for capturing the partial ordering of program statements:
  - Each vertex is the execution of a statement
  - $a \rightarrow b$  means a happens before b
  - Edges can be labeled with current value of variables
  - printf vertices can be labeled with output
  - Each graph begins with a vertex with no incoming edges
- Any topological sort of the graph corresponds to a feasible total ordering (i.e., valid output)
  - A permutation of vertices where all edges point from left to right

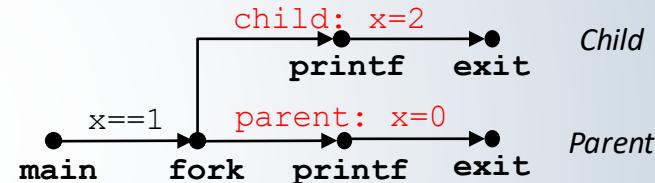
# Process Graph Example

```
int main()
{
    pid_t pid;
    int x = 1;

    pid = fork();
    if (pid == 0) { /* Child */
        printf("child: x=%d\n", ++x);
        exit(0);
    }

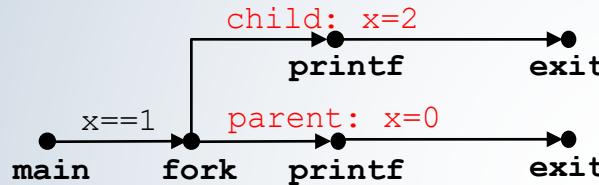
    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}
```

*fork.c*

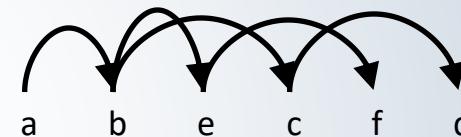


# Interpreting Process Graphs

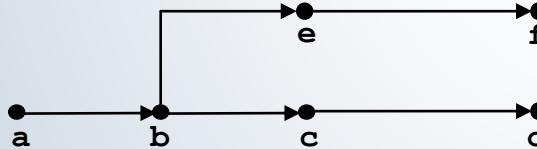
- Original graph:



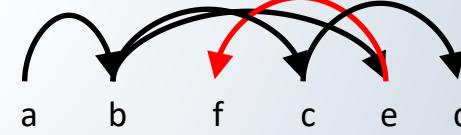
Feasible total ordering:



- Relabeled graph:



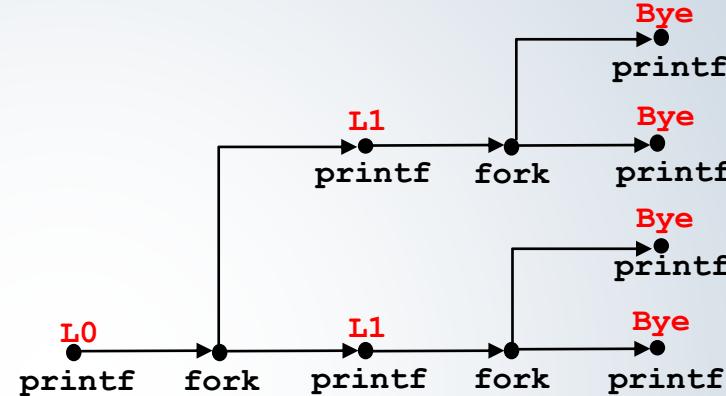
Infeasible total ordering:



# fork() Example: Two Consecutive Forks

```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```

*forks.c*



Feasible output:

L0  
L1  
Bye  
Bye  
L1  
Bye  
Bye  
Bye

Infeasible output:

L0  
Bye  
L1  
Bye  
L1  
Bye  
Bye  
Bye

# fork() Example: Nested Forks in Parent

```
void fork4()
{
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        if (fork() != 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```

*forks.c*

Feasible output:

Infeasible output:

# fork() Example: Nested Forks in Parent

```
void fork4()
{
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        if (fork() != 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```

*forks.c*

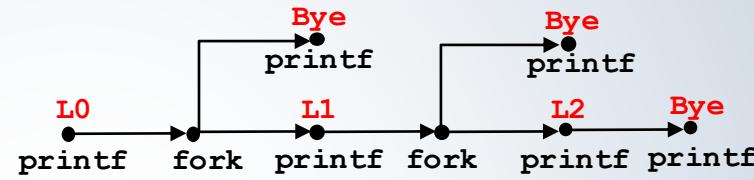
Feasible output:  
L0  
L1  
Bye  
Bye  
L2  
Bye

Infeasible output:  
L0  
Bye  
L1  
Bye  
Bye  
L2

# fork() Example: Nested Forks in Parent

```
void fork4()
{
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        if (fork() != 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```

*forks.c*



Feasible output:

L0  
L1  
Bye  
Bye  
L2  
Bye

Infeasible output:

L0  
Bye  
L1  
Bye  
Bye  
L2

# fork() Example: Nested Forks in Children

```
void fork5()
{
    printf("L0\n");
    if (fork() == 0) {
        printf("L1\n");
        if (fork() == 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```

*forks.c*

Feasible output:

L0  
Bye  
L1  
L2  
Bye  
Bye

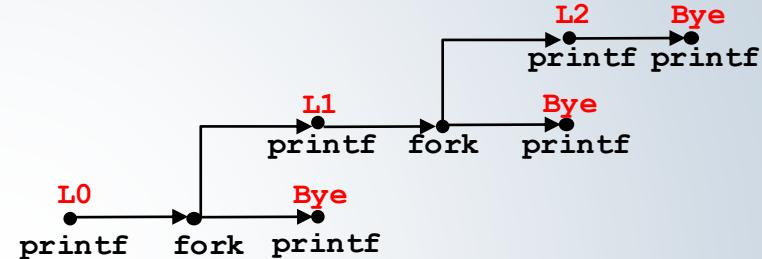
Infeasible output:

L0  
Bye  
L1  
Bye  
Bye  
L2

# fork() Example: Nested Forks in Children

```
void fork5()
{
    printf("L0\n");
    if (fork() == 0) {
        printf("L1\n");
        if (fork() == 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}
```

*forks.c*



Feasible output:

L0  
Bye  
L1  
L2  
Bye  
Bye

Infeasible output:

L0  
Bye  
L1  
Bye  
Bye  
L2

# Reaping Child Processes

- Idea
  - When process terminates, it still consumes system resources
    - Example: Exit status
  - Called a “zombie”
    - Living corpse, half alive and half dead
- Reaping
  - Performed by parent on terminated child (using wait or waitpid)
  - Parent is given exit-status information
  - Kernel then deletes zombie child process
- What if parent doesn't reap?
  - If any parent terminates without reaping a child, then the orphaned child will be reaped by init process (pid == 1)
  - So, only need explicit reaping in long-running processes
    - e.g., shells and servers

# After fork()

- The new process inherits:
  - Process group ID. The child can create a new group and become a leader using setpgid(0, 0)
  - Resource limits
  - Working directory. E.g., filedesc = open("testfile.txt",some options...)
  - Open file descriptors
- But what if we want to execute a different program via fork()?

# exec(): Loading a New Program

- The exec() is a family of functions that load a new program
  - The existing address space is blown away and loaded with the data and instructions of the new program
  - However, things like the PID and file descriptors remain the same
- exec() causes the OS to:
  - Destroy the address space of the calling process
  - Load the new program in memory, creating a new stack and heap
  - Run the new program from its entry point

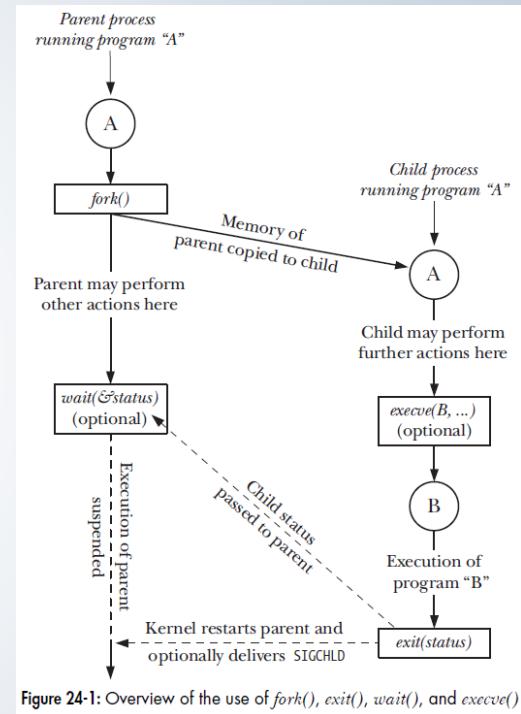


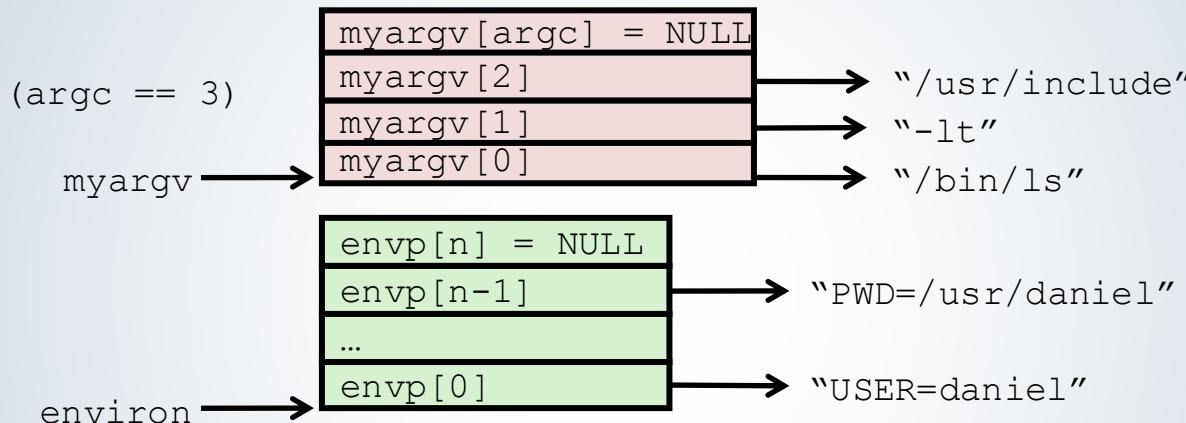
Figure 24-1: Overview of the use of `fork()`, `exit()`, `wait()`, and `execve()`

# execve(): Loading and Running Programs

- `int execve(char *filename, char *argv[], char *envp[])`
- Loads and runs in the current process:
  - Executable file filename
    - Can be object file or script file beginning with `#!/interpreter` (e.g., `#!/bin/bash`)
  - ...with argument list argv
    - By convention `argv[0] == filename`
  - ...and environment variable list envp
    - “name=value” strings (e.g., `USER=droh`)
    - `getenv`, `putenv`, `printenv`
- Overwrites code, data, and stack
  - Retains PID, open files and signal context
- Called once and never returns if it succeeds because `exec()` overwrites the address space of the process
  - ...except if there is an error

# execve() Example

- Executes “/bin/ls -lt /usr/include” in child process using current environment:



```
if ((pid = fork()) == 0) { /* Child runs program */
    if (execve(myargv[0], myargv, environ) < 0) {
        printf("%s: Command not found.\n", myargv[0]);
        exit(1);
    }
}
```

# Zombie Example

```
linux> ./forks 7 &
[1] 6639
Running Parent, PID = 6639
Terminating Child, PID = 6640
```

```
linux> ps
  PID TTY          TIME CMD
 6585 ttys9        00:00:00 tcsh
 6639 ttys9        00:00:03 forks
 6640 ttys9        00:00:00 forks <defunct>
 6641 ttys9        00:00:00 ps
```

```
linux> kill 6639
[1]  Terminated
```

```
linux> ps
  PID TTY          TIME CMD
 6585 ttys9        00:00:00 tcsh
 6642 ttys9        00:00:00 ps
```

```
void fork7() {
    if (fork() == 0) {
        /* Child */
        printf("Terminating Child, PID = %d\n", getpid());
        exit(0);
    } else {
        printf("Running Parent, PID = %d\n", getpid());
        while (1)
            ; /* Infinite loop */
    }
}
```

- ps shows child process as “defunct” (i.e., a zombie)
- Killing parent allows child to be reaped by init. It is known as *cascading* termination. init periodically calls wait() to identify and terminate the orphaned processes

# Non-Terminating Child Example

```
void fork8()
{
    if (fork() == 0) {
        /* Child */
        printf("Running Child, PID = %d\n",
               getpid());
        while (1)
            ; /* Infinite loop */
    } else {
        printf("Terminating Parent, PID = %d\n",
               getpid());
        exit(0);
    }
}
```

```
linux> ./forks 8
Terminating Parent, PID = 6675
Running Child, PID = 6676
linux> ps
      PID TTY          TIME CMD
    6585 tttyp9        00:00:00 tcsh
    6676 tttyp9        00:00:06 forks
    6677 tttyp9        00:00:00 ps
linux> kill 6676
linux> ps
      PID TTY          TIME CMD
    6585 tttyp9        00:00:00 tcsh
    6678 tttyp9        00:00:00 ps
```

- Child process still active even though parent has terminated
- Must kill child explicitly, or else will keep running indefinitely



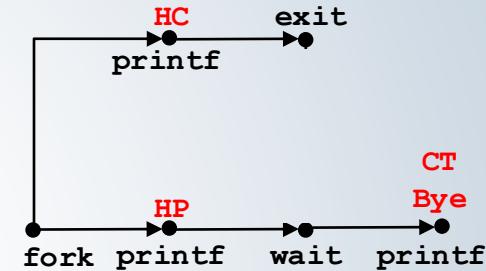
# wait(): Synchronizing with Children

- Parent reaps a child by calling the wait function
- `int wait(int *child_status)`
  - Suspends current process until one of its children terminates
  - Return value is the pid of the child process that terminated
  - `child_status` variable is used to communicate information about the child
    - In Linux, macros defined in `wait.h` (e.g., `WIFEXITED`, `WIFSIGNALED`, etc.) can be used to determine how the process was terminated
      - See textbook for details if interested

# wait(): Synchronizing with Children

```
void fork9() {  
    int child_status;  
  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
        exit(0);  
    } else {  
        printf("HP: hello from parent\n");  
        wait(&child_status);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
}
```

*forks.c*



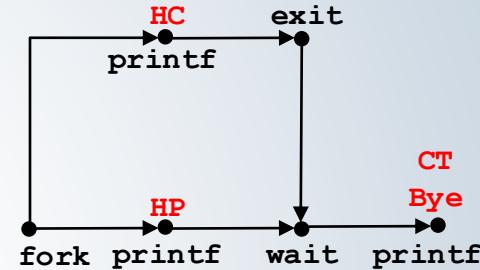
Feasible output:

Infeasible output:

# wait(): Synchronizing with Children

```
void fork9() {  
    int child_status;  
  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
        exit(0);  
    } else {  
        printf("HP: hello from parent\n");  
        wait(&child_status);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
}
```

*forks.c*



Feasible output:

HC  
HP  
CT  
Bye

Infeasible output:

HP  
CT  
Bye  
HC

# Exercise

```
void fork_exercise() {  
    int pid;  
    int status;  
  
    for(i = 0; i < 3; i++){  
        if(i % 2 == 0){  
            pid = fork();  
        }  
        printf("i = %d\n", i);  
    }  
    if(pid != 0){  
        wait(&status);  
    }  
    printf("Bye\n");  
}
```

Draw the process graph

How many times is “Bye” printed?

# Summary

- Spawning processes
  - Call fork()
  - One call, two returns
- Process completion
  - Call exit()
  - One call, no return
- Reaping and waiting for processes
  - Call wait() or waitpid()
- Loading and running programs
  - Call execve() (or variant)
  - One call, no return unless error