

CS3281 / CS5281

Virtual Memory

CS3281 / CS5281

Spring 2026

**Some lecture slides borrowed and adapted from CMU's
“Computer Systems: A Programmer’s Perspective”*


Today

- Address spaces
- VM as a tool for memory management
- VM as a tool for caching
- VM as a tool for memory protection
- Address translation




A System Using Physical Addressing

- Used in “simple” systems like embedded microcontrollers in devices like cars, elevators, and digital picture frames



A System Using Virtual Addressing

- Used in all modern servers, laptops, and smart phones
- One of the great ideas in computer science



Address Spaces

- **Linear address space:** Ordered set of contiguous non-negative integer addresses:
 $\{0, 1, 2, 3 \dots\}$
- **Physical address space:** Set of $M = 2^m$ physical addresses
 $\{0, 1, 2, 3, \dots, M-1\}$
 - If you have 16 GB of RAM, you need at least 34 bits of byte addressing ($2^{34} = 16 \text{ GiB}$) to address all bytes.
- **Virtual address space:** Set of $N = 2^n$ virtual addresses
 $\{0, 1, 2, 3, \dots, N-1\}$
 - If a machine uses **48-bit virtual addresses**, then the virtual address space size is 2^{48} bytes.

Why Virtual Memory (VM)?

- Simplifies memory management
 - Each process gets the same linear address space
- Isolates address spaces
 - One process can't interfere with another's memory
- Uses main memory (RAM) efficiently
 - Use DRAM as a cache for parts of a virtual address space. A very large virtual memory is perceived by programmers when only a small physical memory is available
- Many other benefits (some discussed later)
 - Shared memory, memory deduplication, lazy allocation, etc.


Today

- Address spaces
- VM as a tool for memory management
- VM as a tool for caching
- VM as a tool for memory protection
- Address translation




VM as a Tool for Memory Management

- Key idea: each process has its own virtual address space
 - It can view memory as a simple linear array
 - Mapping function scatters addresses through physical memory
 - Page: a fixed-size block of memory (e.g., 4KB)



VM as a Tool for Memory Management

- **Simplifying memory allocation**
 - Each virtual page can be mapped to any physical page
 - A virtual page can be stored in different physical pages at different times
- **Sharing code and data among processes**
 - Map virtual pages to the same physical page (here: PP 6)




Today

- Address spaces
- VM as a tool for memory management
- VM as a tool for caching
- VM as a tool for memory protection
- Address translation




Latencies in Different Levels of Memories

Latency Numbers Every Programmer Should Know



VM as a Tool for Caching

- Conceptually, *virtual memory* is an array of N contiguous bytes stored on disk.
- The contents of the array on disk are cached in *physical memory (DRAM cache)*
 - These cache blocks are called *pages* (size is $P = 2^p$ bytes). Pages may (not) be contiguous in the physical memory



DRAM Cache Organization


- DRAM cache organization driven by the enormous miss penalty
 - DRAM is about **10x** slower than SRAM (cache)
 - Disk is about **10,000x** slower than DRAM
- Consequences (page fault)
 - Large page (block) size: typically 4 KB (Huge pages are 2MB – 1GB.)
 - Only a subset of virtual pages are stored in the main memory (working set)
 - Highly sophisticated, expensive replacement algorithms
 - Too complicated and open-ended to be implemented in hardware

Internal Fragmentation

- Depending on a page size, a program size may not be a multiple of the number of pages. Thus, the last page is partially filled. This loss of usable memory is known as internal fragmentation
- The OS prevents out-of-range accesses by leaving pages outside the allocated address range unmapped; accesses to unmapped pages fault (segfault). Internal fragmentation remains because the last page may be only partially used.


Enabling Data Structure: Page Table

- A *page table* is an array of page table entries (PTEs) that maps virtual pages to physical pages.
 - Per-process kernel data structure in DRAM




Page Hit

- **Page hit:** reference to VM word that is in physical memory (DRAM cache hit)




Page Fault

- **Page fault:** reference to VM word not in physical memory (DRAM cache miss). The VM word can be instruction or data




Handling Page Fault

- Page miss causes page fault (an exception)




Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)




Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)




Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)
- Offending (faulting) instruction is restarted: page hit! If the page fault is caused by memory miss, the data operand is fetched from memory



Allocating Pages




`malloc()` allocates a new page (VP 5) of virtual memory. This page stays in disk until it is used

Today

- Address spaces
- VM as a tool for memory management
- VM as a tool for caching
- VM as a tool for memory protection
- Address translation

VM as a Tool for Memory Protection

- Extend PTEs with permission bits
- MMU checks these bits on each access



Today

- Address spaces
- VM as a tool for memory management
- VM as a tool for memory protection
- VM as a tool for caching
- Address translation




VM Address Translation

- Virtual Address Space
 - $V = \{0, 1, \dots, N-1\}$
- Physical Address Space
 - $P = \{0, 1, \dots, M-1\}$
- Address Translation
 - **MAP:** $V \rightarrow P \cup \{\emptyset\}$
 - For virtual address a :
 - $MAP(a) = a'$ if data at virtual address a is at physical address a' in P
 - $MAP(a) = \emptyset$ if data at virtual address a is not in physical memory (page fault)


Summary of Address Translation Symbols

- Basic Parameters
 - **n, m, p**: Number of bits
 - **N = 2^n** : Number of addresses in virtual address space
 - **M = 2^m** : Number of addresses in physical address space
 - **P = 2^p** : Page size (bytes)
- Components of the virtual address (VA)
 - **VPO**: Virtual page offset
 - **VPN**: Virtual page number
- Components of the physical address (PA)
 - **PPO**: Physical page offset (same as VPO)
 - **PPN**: Physical page number

Address Translation with a Page Table




Address Translation: Page Hit



- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) MMU sends physical address to cache/memory
- 5) Cache/memory sends data word to processor


Address Translation: Page Fault




- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) Valid bit is zero, so MMU triggers page fault exception
- 5) Handler identifies victim (and, if dirty, pages it out to disk)
- 6) Handler pages in new page and updates PTE in memory
- 7) Handler returns to original process, restarting faulting instruction

Multi-Level Page Tables


- Suppose:
 - 4KB (2^{12}) page size, 48-bit address space, 8-byte PTE (2^3)
- Problem:
 - Would need a 512 GB page table!
 - $2^{48} * 2^{-12} * 2^3 = 2^{39}$ bytes
- Common solution: Multi-level page table
- Example: 2-level page table
 - Level 1 table: each PTE points to a page table (always memory resident)
 - Level 2 table: each PTE points to a data page (paged in and out like any other data)



Two-Level Page-Table Hierarchy



Translating with a 2-Level Page Table



Multi-level page table (example)

- Consider a 3-level page table. A page table has 64 entries. If the page size is 4KB, how much is the size of the virtual address space?

Summary

- Programmer's view of virtual memory
 - Each process has its own private linear address space
 - Cannot be corrupted by other processes
- System view of virtual memory
 - Uses memory efficiently by caching virtual memory pages
 - Efficient only because of locality
 - Simplifies memory management and programming
 - Simplifies protection by providing a convenient point to check permissions