



CS3281 / CS5281

Exceptional Control Flow

CS3281 / CS5281

Spring 2024

**Some lecture slides borrowed and adapted from CMU's
"Computer Systems: A Programmer's Perspective"*



Tel (615) 343-7472 | Fax (615) 343-7440
1025 16th Avenue South Nashville, TN 37212
www.isis.vanderbilt.edu



This Lecture

- **Exceptional Control Flow**
- Exceptions
- Processes
- Process Control



Control Flow

- Processors do only one thing:
 - From startup to shutdown, a CPU reads and executes (interprets) a sequence of instructions, one at a time
 - This sequence is the CPU's *control flow* (or *flow of control*)

Time



Physical control flow

<startup>

inst₁

inst₂

inst₃

...

inst_n

<shutdown>

Altering the Control Flow

- You know two mechanisms for changing control flow:
 - Jumps and branches
 - Call and return
- Insufficient for a useful system: difficult to react to events and changing system state
 - Examples of changes in system state:
 - Data arrives from a disk or a network adapter
 - Instruction divides by zero
 - User hits Ctrl-C at the keyboard
 - System timer expires
- System needs mechanisms for “exceptional control flow”

Exceptional Control Flow

- Exists at all levels of a computer system
- Low-level mechanisms:
 - **Exceptions:** change in control flow in response to a system event
 - Implemented using combination of hardware and OS software
- High-level mechanisms:
 - **Process context switch**, i.e., stop running one program and start running another
 - Implemented by OS software and hardware timer
 - **Signals:** a means of sending a (limited) message to a process that they should respond to
 - Implemented by OS software
 - **Nonlocal jumps:** `setjmp()` and `longjmp()`
 - Implemented by C runtime library

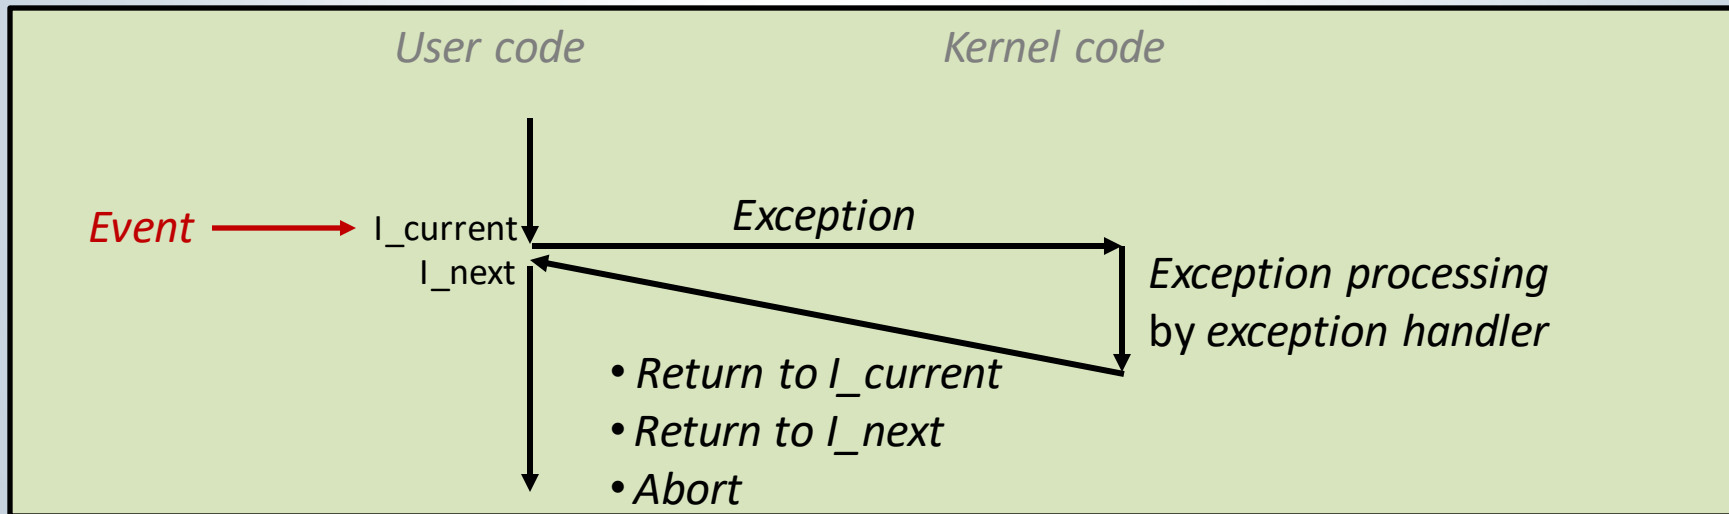
This Lecture

- Exceptional Control Flow
- **Exceptions**
- Processes
- Process Control

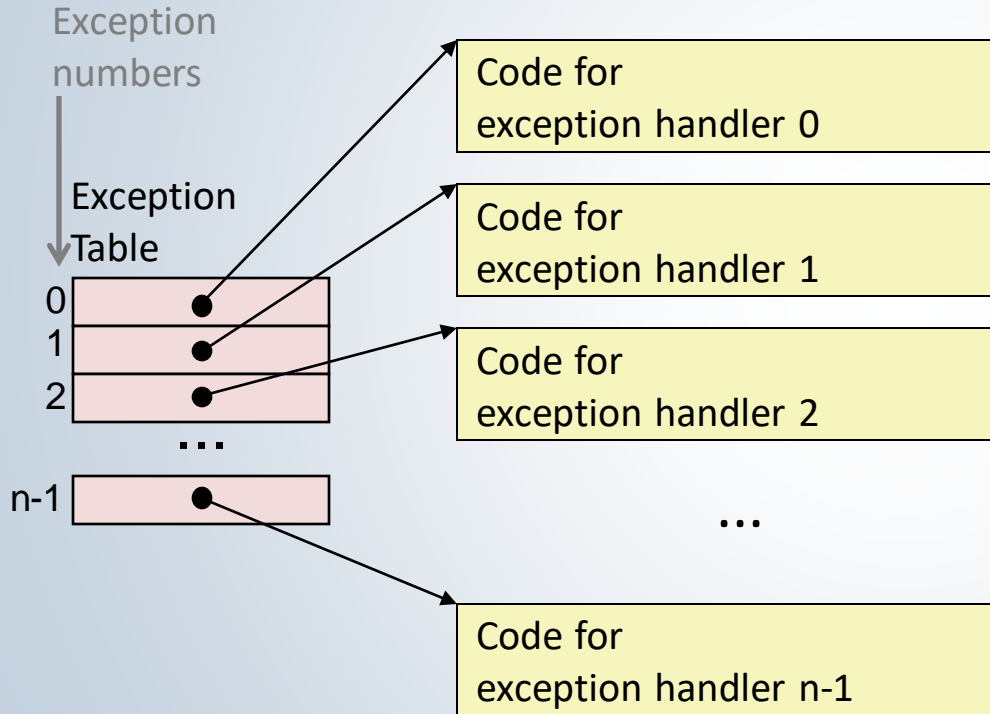


Exceptions

- An exception is a transfer of control to the OS kernel in response to some event (i.e., change in processor state)
 - Examples of events: Divide by 0, arithmetic overflow, page fault, I/O request completes, typing Ctrl-C



Exception Tables



- Each type of event has a unique exception number k
- k = index into exception table (a.k.a. interrupt vector)
- Handler k is called each time exception k occurs

Asynchronous Exceptions

- Asynchronous exceptions are called interrupts
- Caused by events external to the processor
 - Indicated by setting the processor's interrupt pin
 - Handler returns to “next” instruction
- Examples:
 - Timer interrupt
 - Every few ms, a timer chip trigger an interrupt
 - Used by kernel to take back control from user programs
 - I/O interrupt from external device
 - Arrival of network packet
 - Arrival of data from disk

Synchronous Exceptions

- Caused by events that occur as a result of executing an instruction:
 - Traps
 - Intentional
 - Examples: syscalls, breakpoint traps, special instructions
 - Returns control to “next” instruction
 - Faults
 - Unintentional but possibly recoverable
 - Examples: page faults (recoverable), protection faults (unrecoverable), floating-point exceptions
 - Either re-execute faulting instruction or abort
 - Aborts
 - Unintentional and unrecoverable
 - Examples: illegal instruction, parity error, machine check
 - Aborts current program

System Calls


- Each system call has a unique ID number
- xv6 syscall numbers defined in kernel/syscall.h
- Linux has many more system calls as it is much more complex
- You will be implementing your own system calls, so you will have to add to this

```
// System call numbers
#define SYS_fork      1
#define SYS_exit      2
#define SYS_wait      3
#define SYS_pipe      4
#define SYS_read      5
#define SYS_kill      6
#define SYS_exec      7
#define SYS_fstat     8
#define SYS_chdir     9
#define SYS_dup      10
#define SYS_getpid    11
#define SYS_sbrk     12
#define SYS_sleep     13
#define SYS_uptime    14
#define SYS_open      15
#define SYS_write     16
#define SYS_mknod     17
#define SYS_unlink    18
#define SYS_link      19
#define SYS_mkdir     20
#define SYS_close     21
```

System Call Example: Opening File

- User calls: `open(filename, options)`
- Calls `__open` function, which invokes system call instruction `syscall`

```
0000000000e5d70 <__open>:  
...  
e5d79: b8 02 00 00 00    mov $0x2,%eax      # open is syscall #2  
e5d7e: 0f 05             syscall            # Return value in %rax  
e5d80: 48 3d 01 f0 ff ff  cmp $0xfffffffffff001,%rax  
...  
e5dfa: c3               retq
```



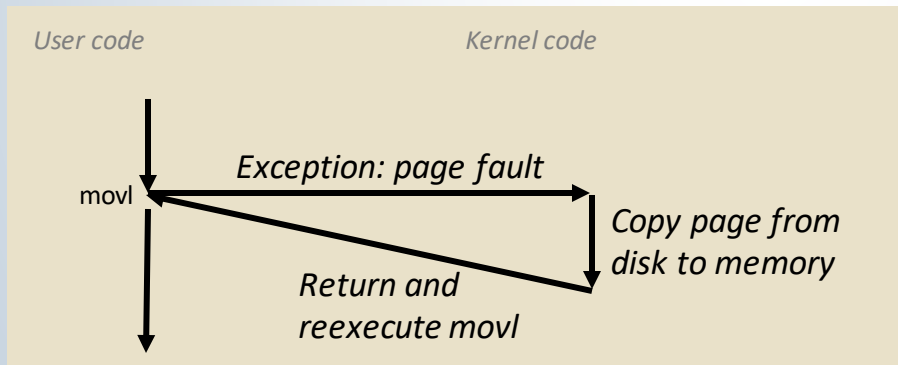
- `%eax` contains syscall number
- Other arguments in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`
- Return value in `%rax`
- Negative value is an error corresponding to negative `errno`

Fault Example: Page Fault

- User writes to memory location
- That portion (page) of user's memory is currently on disk

```
80483b7: c7 05 10 9d 04 08 0d    movl    $0xd,0x8049d10
```

```
int a[1000];  
main ()  
{  
    a[500] = 13;  
}
```

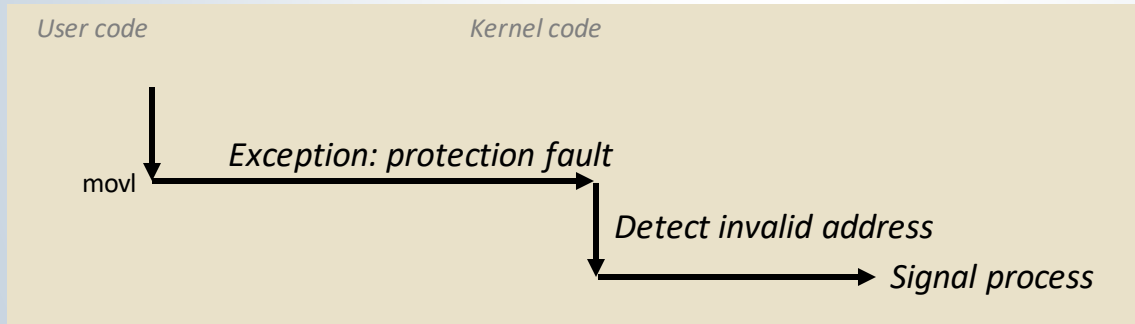


Fault Example: Invalid Memory Reference

- Buffer overflow
- Sends `SIGSEGV` signal to user process
- User process exits with “segmentation fault”

```
int a[1000];  
main ()  
{  
    a[5000] = 13;  
}
```

```
80483b7: c7 05 60 e3 04 08 0d    movl    $0xd,0x804e360
```




This Lecture

- Exceptional Control Flow
- Exceptions
- **Processes**
- Process Control



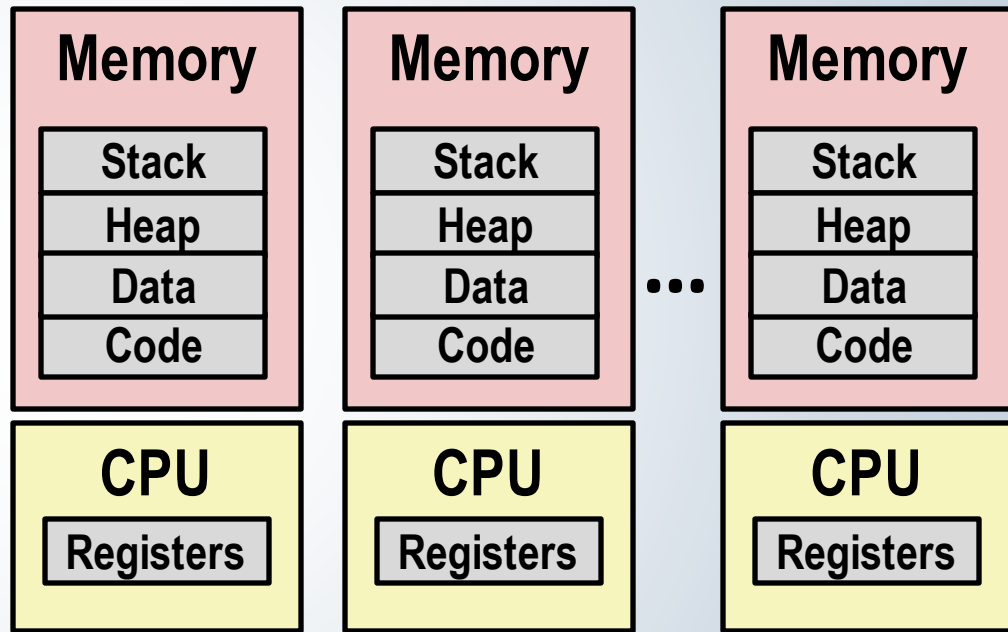
Processes

- Definition: A process is an instance of a running program
 - One of the most profound ideas in computer science
 - Not the same as “program” or “processor”
- Process provides each program with two key abstractions:
 - Logical control flow
 - Each program seems to have exclusive use of the CPU
 - Provided by kernel mechanism called context switching
 - Private address space
 - Each program seems to have exclusive use of main memory
 - Provided by kernel mechanism called virtual memory



Multiprocessing: The Illusion

- Computer runs many processes simultaneously
 - Applications for one or more users
 - Web browsers, email clients, editors, ...
 - Background tasks
 - Monitoring network & I/O devices



Multiprocessing Example

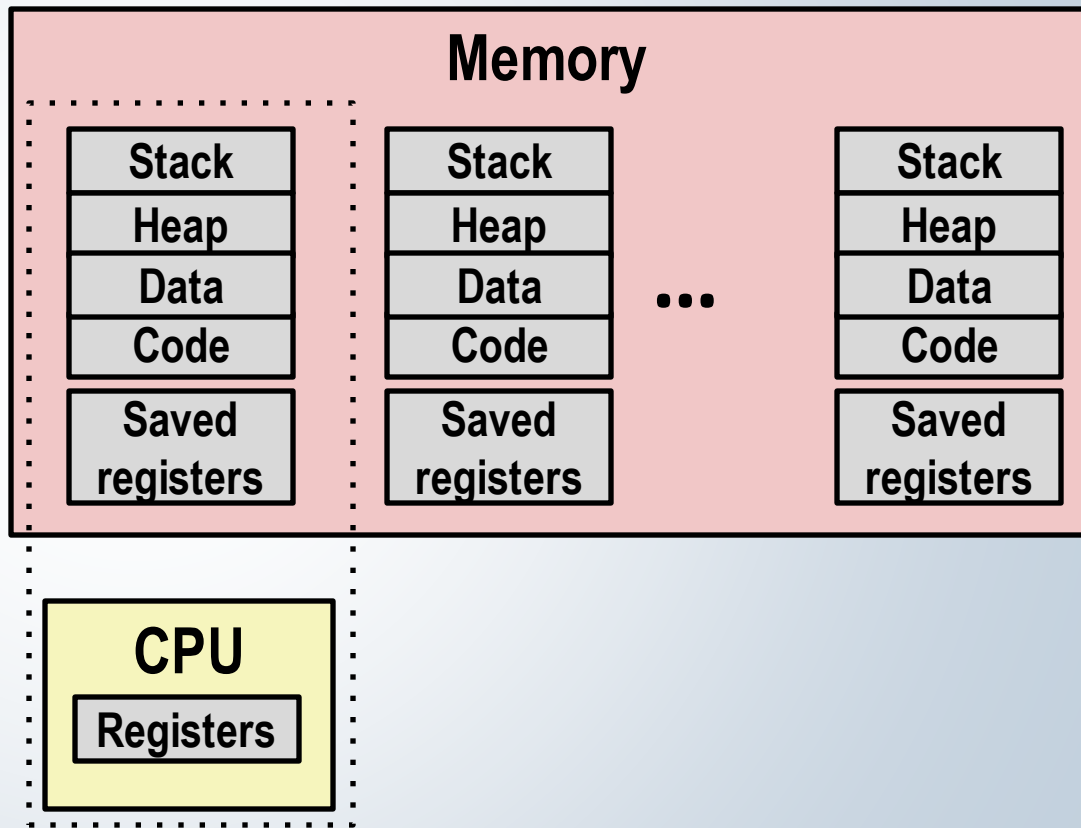
- Running program “top” on Mac
 - System has 123 processes, 5 of which are active
 - Identified by Process ID (PID)

```
xterm
Processes: 123 total, 5 running, 9 stuck, 109 sleeping, 611 threads
Load Avg: 1.03, 1.13, 1.14 CPU usage: 3.27% user, 5.15% sys, 91.56% idle
SharedLibs: 576K resident, 0B data, 0B linkedit.
MemRegions: 27958 total, 1127M resident, 35M private, 494M shared.
PhysMem: 1039M wired, 1974M active, 1062M inactive, 4076M used, 18M free.
VM: 280G vsize, 1091M framework vsize, 23075213(1) pageins, 5843367(0) pageouts.
Networks: packets: 41046228/11G in, 66083096/77G out.
Disks: 17874391/349G read, 12847373/594G written.

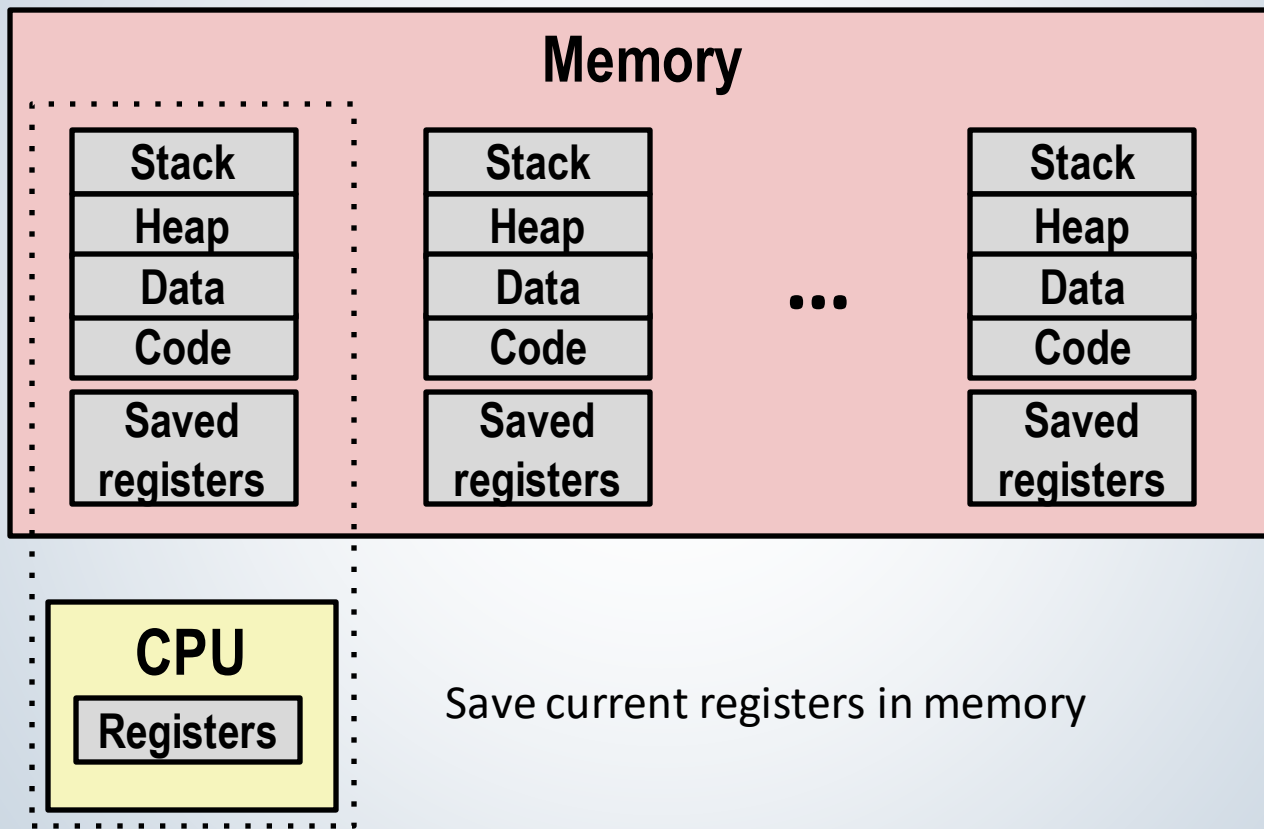
PID    COMMAND    %CPU TIME    #TH    #WQ    #PORT    #HREG    RPRVT    RSHRD    RSIZE    VPRVT    VSIZE
99217-  Microsoft  0.0 02:28.34 4      1      202     418     21M     24M     21M     66M     763M
99051   usbmuxd    0.0 00:04.10 3      1      47      66     436K    216K    480K    60M     2422M
99006   iTunesHelper 0.0 00:01.23 2      1      55      78     728K    3124K   1124K   43M     2429M
84286   bash       0.0 00:00.11 1      0      20      24     224K    732K    484K    17M     2378M
84285   xterm      0.0 00:00.83 1      0      32      73     656K    872K    692K    9728K   2382M
55939-  Microsoft  0.3 21:58.97 10     3      360     954     16M     65M     46M     114M    1057M
54751   sleep      0.0 00:00.00 1      0      17      20     92K     212K    360K    9632K   2370M
54739   launchdadd 0.0 00:00.00 2      1      33      50     488K    220K    1736K   48M     2409M
54737   top        6.5 00:02.53 1/1    0      30      29     1416K   216K    2124K   17M     2378M
54719   automountd 0.0 00:00.02 7      1      53      64     860K    216K    2184K   53M     2413M
54701   ocspd      0.0 00:00.05 4      1      61      54     1268K   2644K   3132K   50M     2426M
54661   Grab       0.6 00:02.75 6      3      222+    389+    15M+    26M+    40M+    75M+    2556M+
54659   cookied    0.0 00:00.15 2      1      40      61     3316K   224K    4088K   42M     2411M
53818   mdworker   0.0 00:01.67 4      1      52      91     7628K   7412K   16M     48M     2438M
50878   mdworker   0.0 00:11.17 3      1      53      91     2464K   6148K   9976K   44M     2434M
50410   xterm      0.0 00:00.13 1      0      32      73     280K    872K    532K    9700K   2382M
50078   emacs      0.0 00:06.70 1      0      20      35     52K     216K    88K     18M     2392M
```

Multiprocessing: The (Traditional) Reality

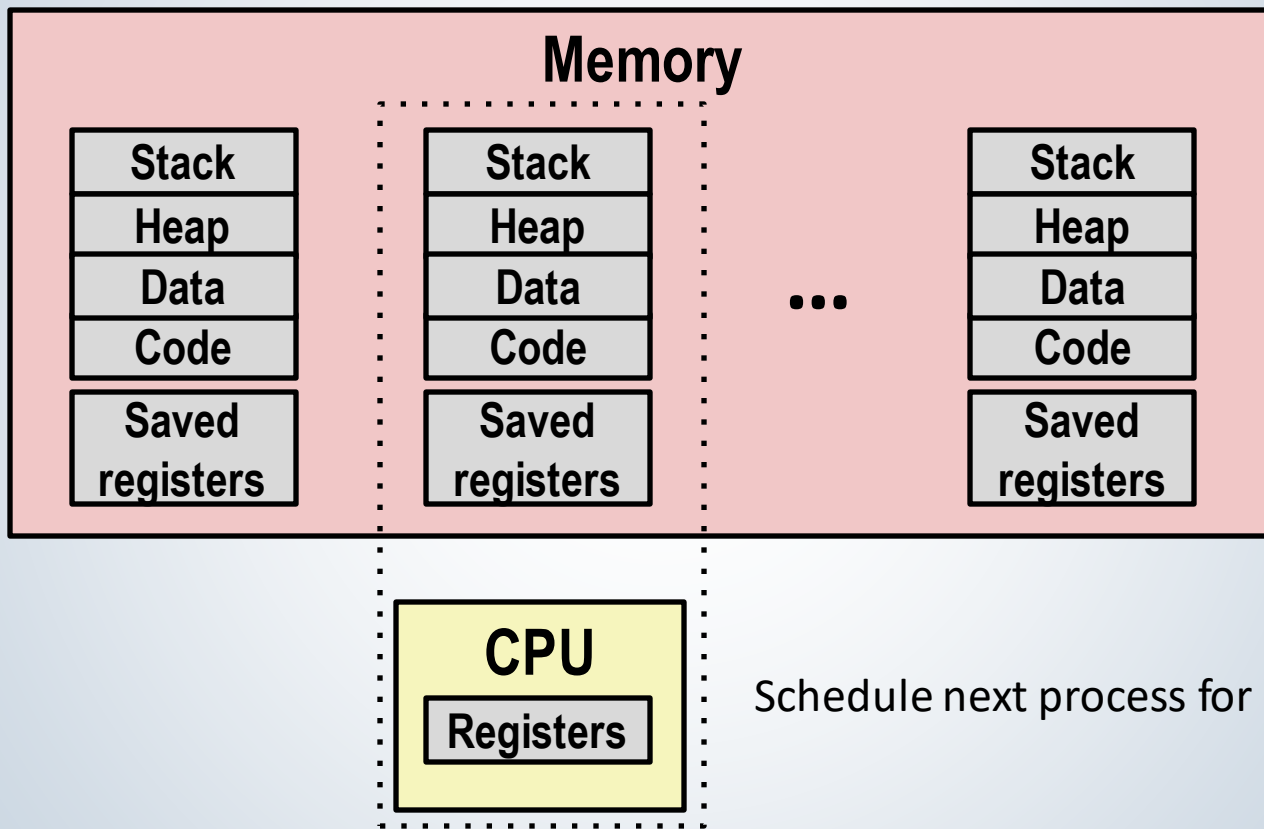
- Single processor executes multiple processes concurrently
 - Process executions interleaved (multitasking)
 - Address spaces managed by virtual memory system (later in course)
 - Register values for nonexecuting processes saved in memory



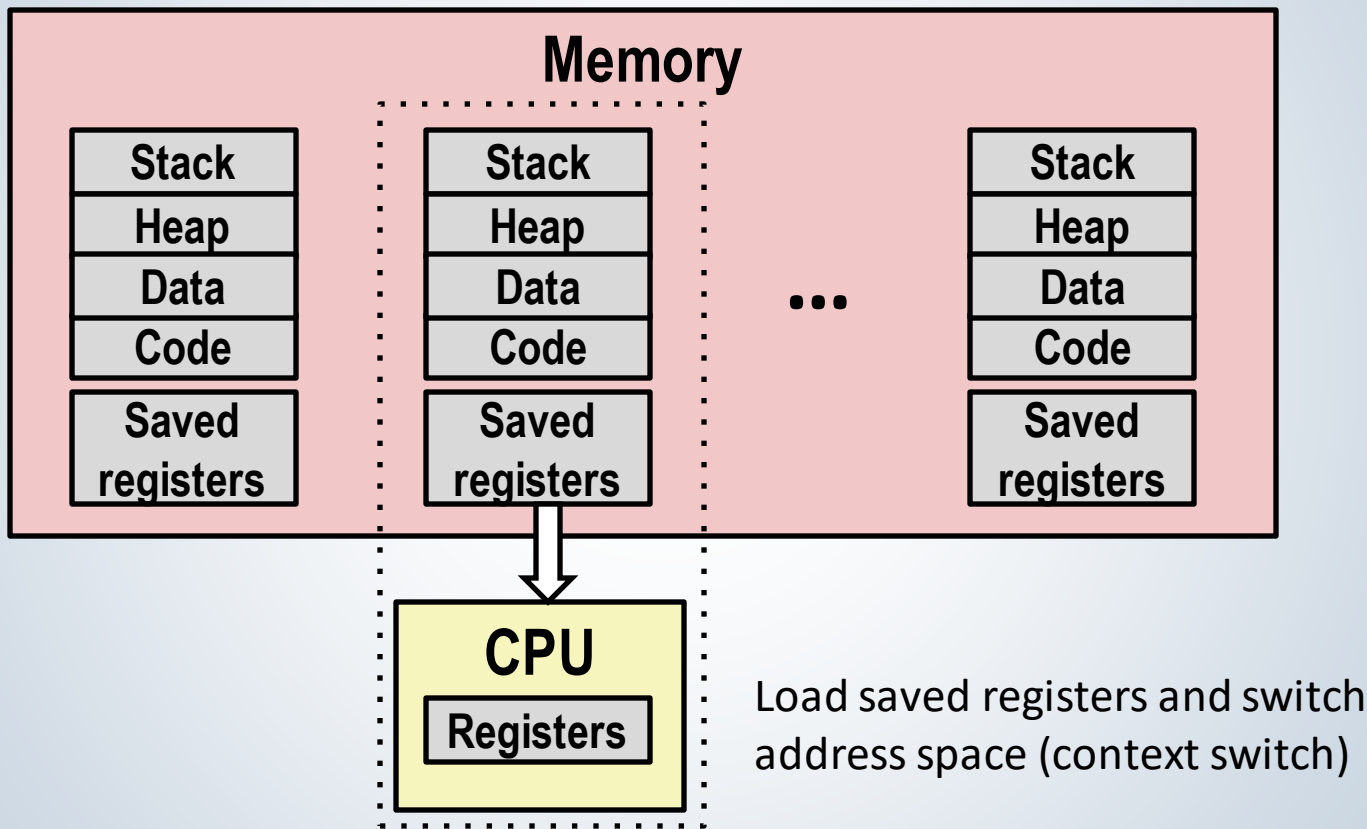
Multiprocessing: The (Traditional) Reality



Multiprocessing: The (Traditional) Reality



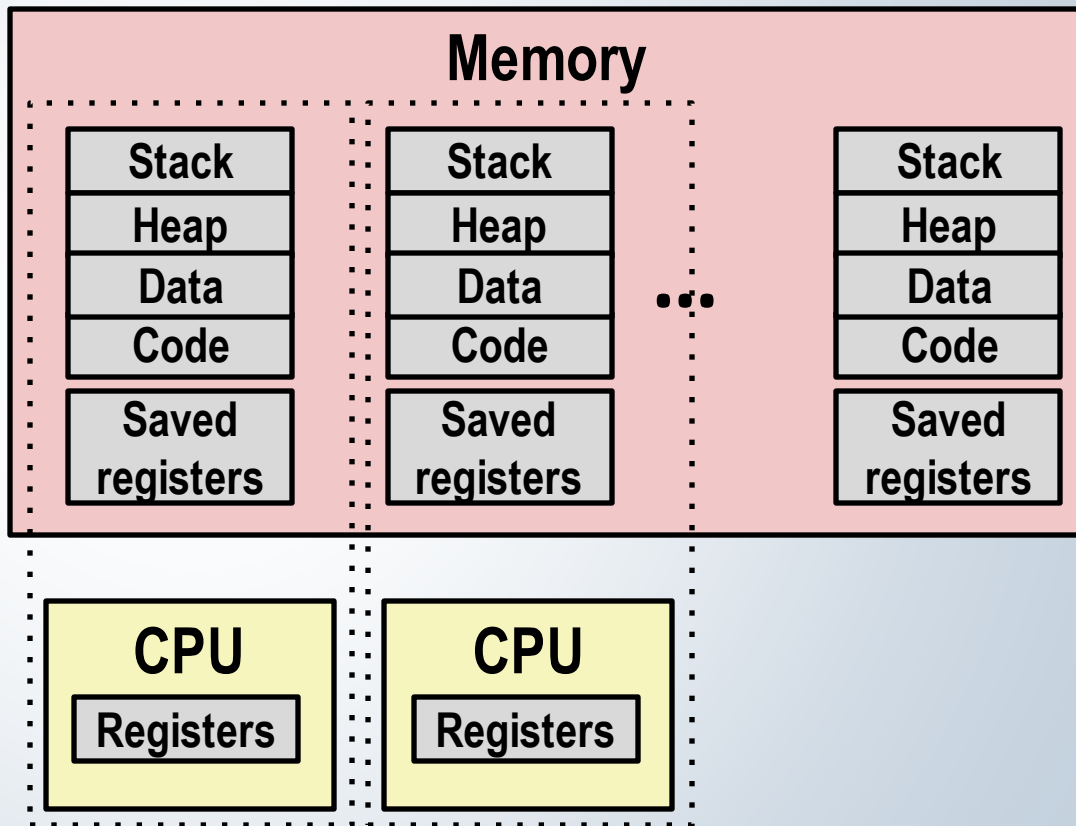
Multiprocessing: The (Traditional) Reality



Multiprocessing: The (Modern) Reality


- **Multicore processors**

- Multiple CPUs on single chip
- Share main memory (and some of the caches)
- Each can execute a separate process
- Scheduling of processors onto cores done by kernel




Concurrent Processes

- Each process is a logical control flow
- Two processes *run concurrently* (are *concurrent*) if their flows overlap in time
- Otherwise, they are *sequential*
- Examples (running on single core):
 - Concurrent: A & B, A & C
 - Sequential: B & C



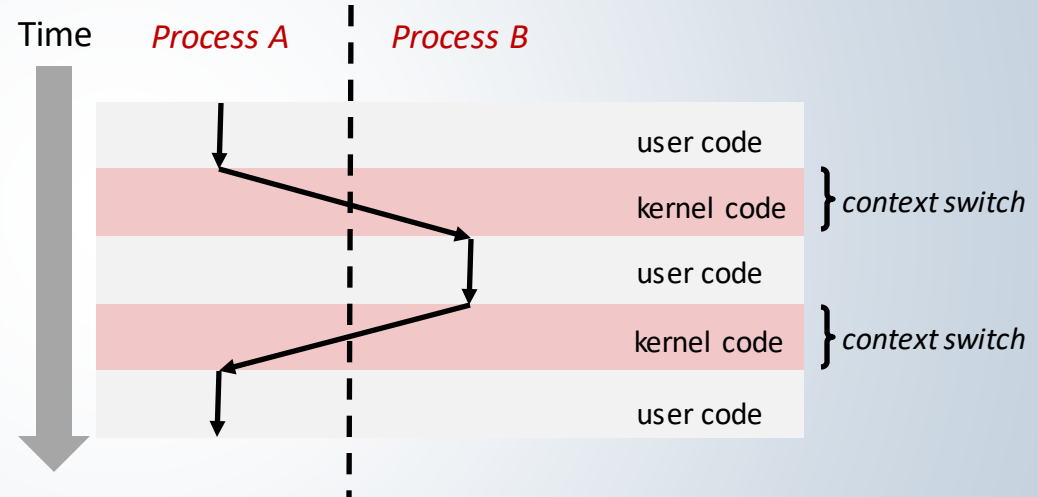
User View of Concurrent Processes

- Control flows for concurrent processes are physically disjoint in time
- However, we can think of concurrent processes as running in parallel with each other



Context Switching

- Processes are managed by a shared chunk of memory-resident OS code called the *kernel*
 - Important:** the kernel is not a separate process, but rather runs as part of some existing process
- Control flow passes from one process to another via a *context switch*



Summary

- Exceptions
 - Events that require nonstandard control flow
 - Generated externally (interrupts) or internally (traps and faults)
- Processes
 - At any given time, system has multiple active processes
 - Only one can execute at a time on a single core, though
 - Each process appears to have total control of processor + private memory space