



CS3281 / CS5281

Virtual Memory (contd.)

CS3281 / CS5281

Fall 2025

**Some lecture slides borrowed and adapted from CMU's "Computer Systems: A Programmer's Perspective", Ghena, St-Amour, Hardavellas, Bustamante (Northwestern), Bryant, O'Hallaron (CMU), Garcia, Weaver (UC Berkeley)*

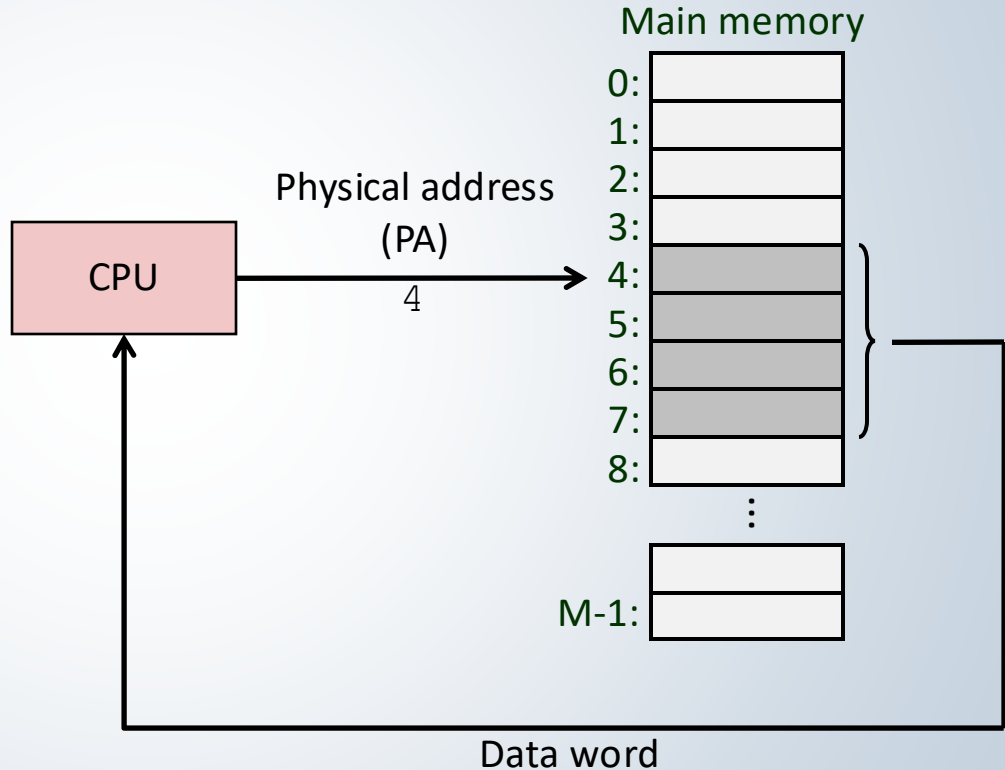


Tel (615) 343-7472 | Fax (615) 343-7440
1025 16th Avenue South Nashville, TN 37212
www.isis.vanderbilt.edu



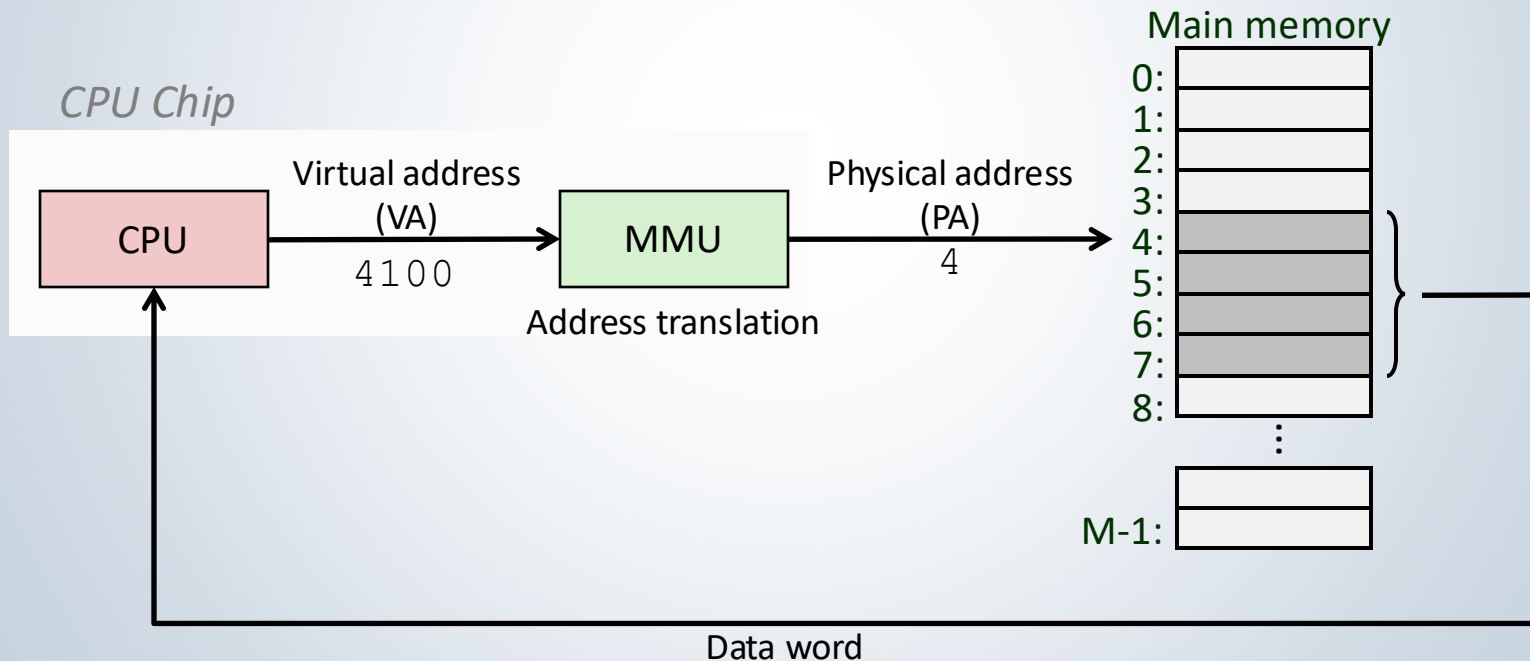
A System Using Physical Addressing

- Used in “simple” systems like embedded microcontrollers in devices like cars, elevators, and digital picture frames



A System Using Virtual Addressing

- Used in all modern servers, laptops, and smart phones
- One of the great ideas in computer science




Address Spaces

- **Linear address space:** Ordered set of contiguous non-negative integer addresses:
 $\{0, 1, 2, 3 \dots \}$
- Number of bits for virtual/physical addressing tell us how big they are

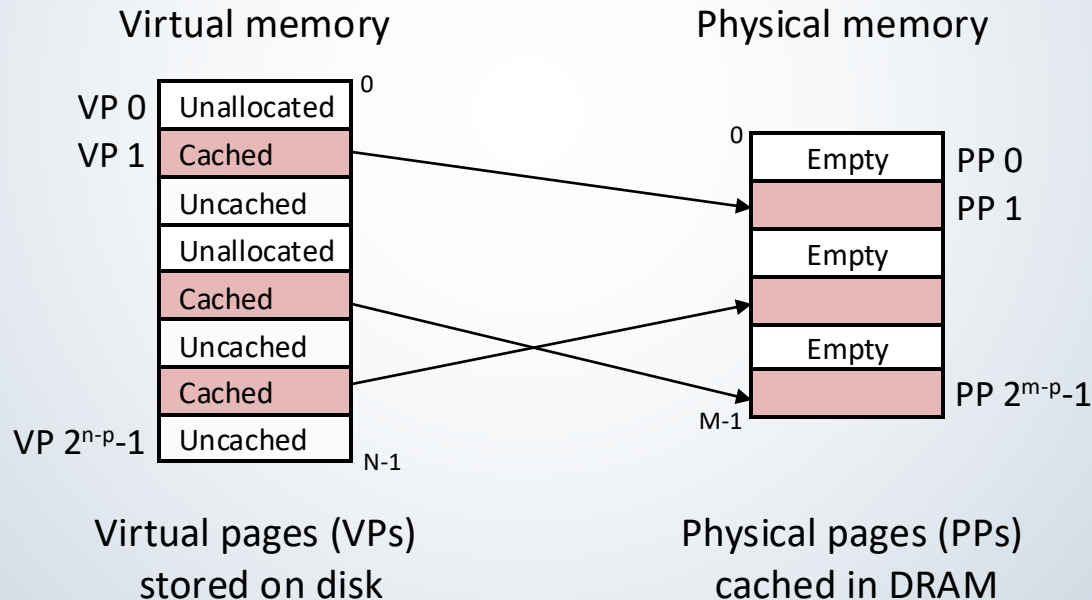
VM as a Tool for Memory Management

- **Key idea: each process has its own virtual address space**
 - It can view memory as a simple linear array
 - Mapping function scatters addresses through physical memory



VM as a Tool for Caching

- Conceptually, *virtual memory* is an array of N contiguous bytes stored on disk.
- The contents of the array on disk are cached in *physical memory (DRAM cache)*
 - These cache blocks are called *pages* (size is $P = 2^p$ bytes). Pages may (not) be contiguous in the physical memory



DRAM Cache Organization

- DRAM cache organization driven by the enormous miss penalty
 - DRAM is about **10x** slower than SRAM (cache)
 - Disk is about **10,000x** slower than DRAM
- Consequences
 - Only a subset of virtual pages are stored in the main memory
 - Highly sophisticated, expensive replacement algorithms
 - Large page (block) size: typically 4 KB
 - CPU caches (SRAM) move data at cache-line granularity (aka “cache blocks”), e.g., 64 B, between SRAM \leftrightarrow DRAM.
 - Virtual memory moves data at page granularity, e.g., 4 KB, between DRAM \leftrightarrow disk. Because the miss penalty to disk is enormous, the transfer unit is much larger than a cache line.

Why pages help: Spatial Locality

- Programs tend to touch nearby addresses soon after one another (arrays, stack frames, sequential code).
- A page fault brings in a whole page (≈ 4 KB), not just one word \rightarrow the next few KB of nearby data/code are already in DRAM.
 - This amortizes the enormous disk miss penalty (disk \gggg DRAM), turning one slow I/O into many fast DRAM hits.
- Typical wins:
 - Instruction streams (sequential fetch within functions)
 - Stacks (new frames/locals)
 - Arrays & buffers (for $i = 0 \dots N$ loops)
 - Contiguous allocations

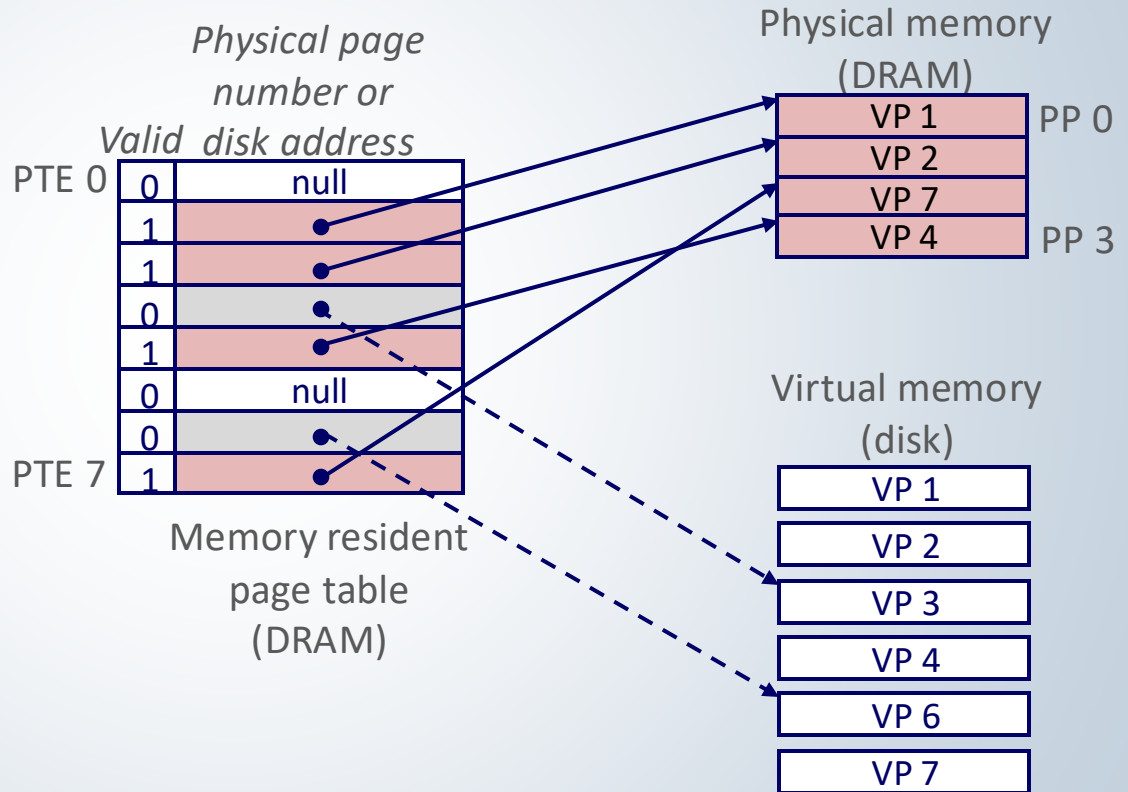
*Works **only if** locality exists; random sparse access wastes I/O and can cause thrashing.*

Internal Fragmentation

- Depending on a page size, a program size may not be a multiple of the number of pages. Thus, the last page is partially filled. This loss of usable memory is known as internal fragmentation
- malloc() usually tries to make unused spaces available

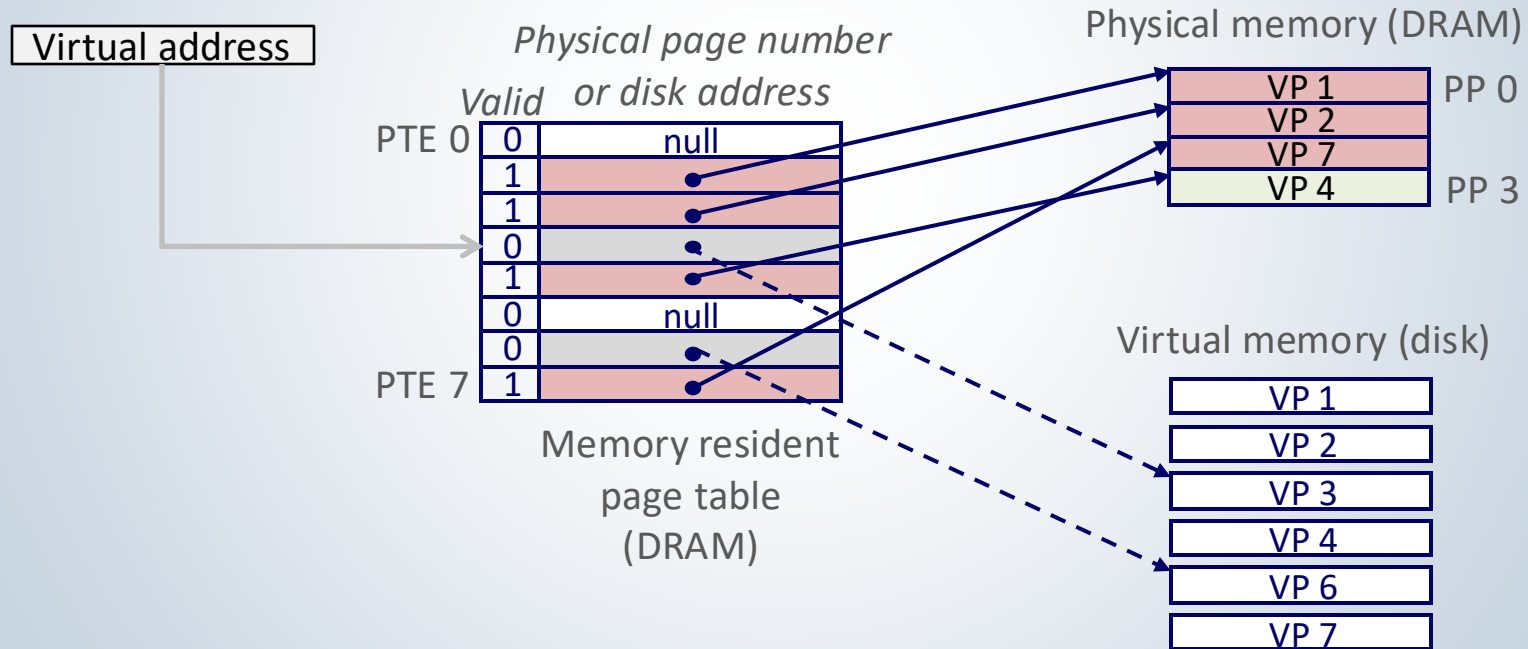
Enabling Data Structure: Page Table

- A *page table* is an array of page table entries (PTEs) that maps virtual pages to physical pages.
 - Per-process kernel data structure in DRAM



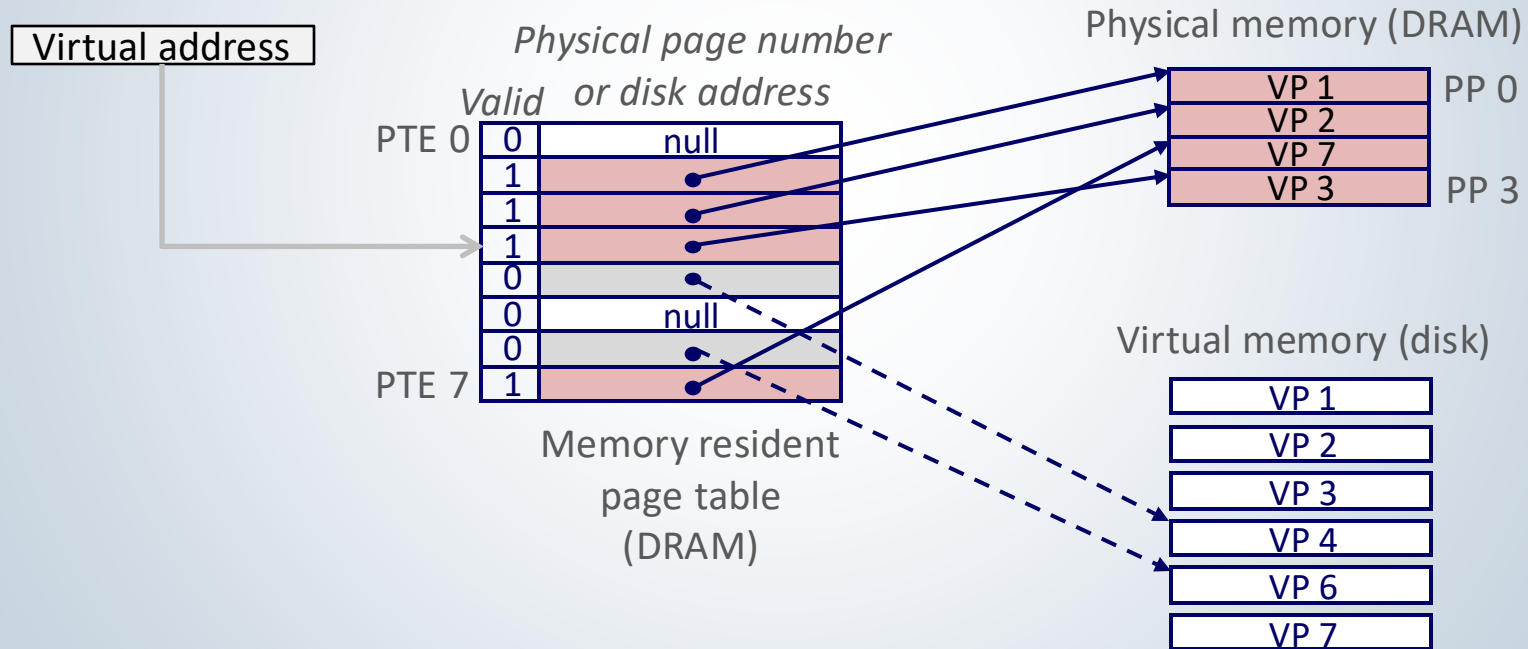
Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)



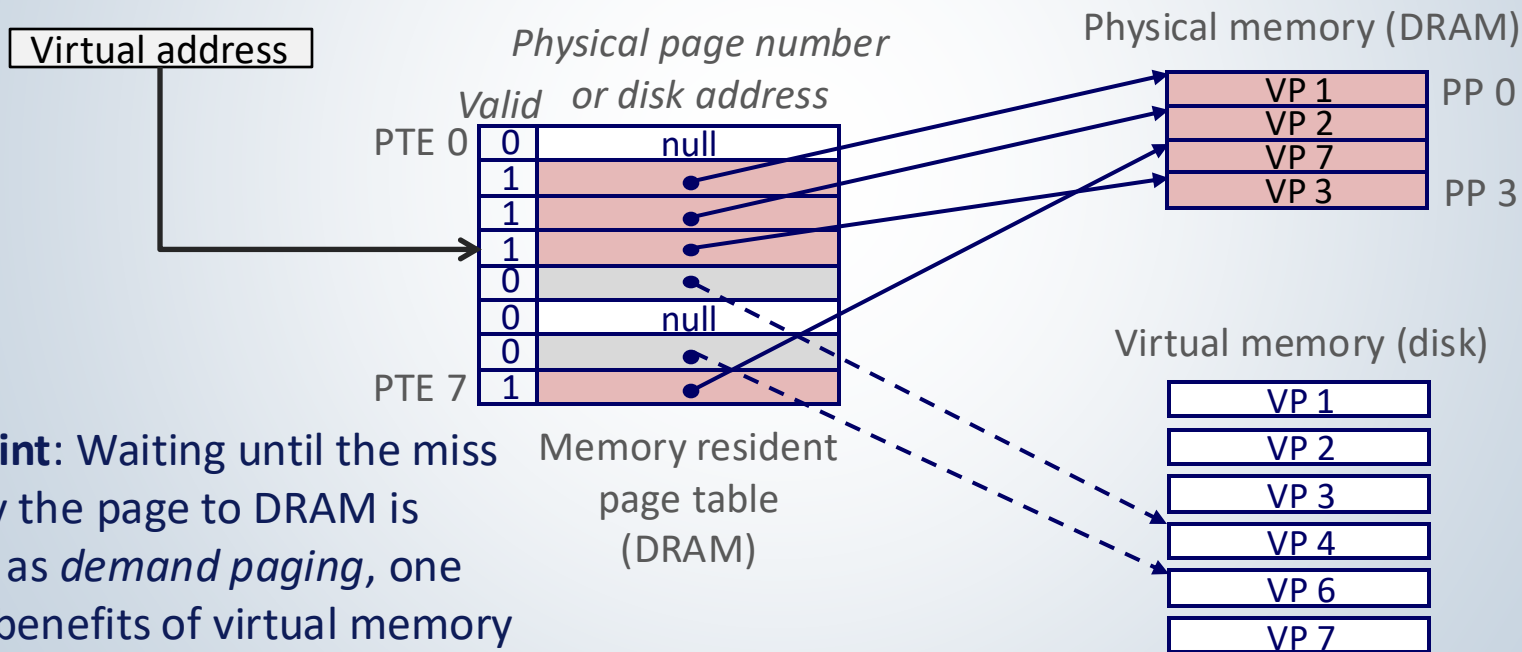
Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)

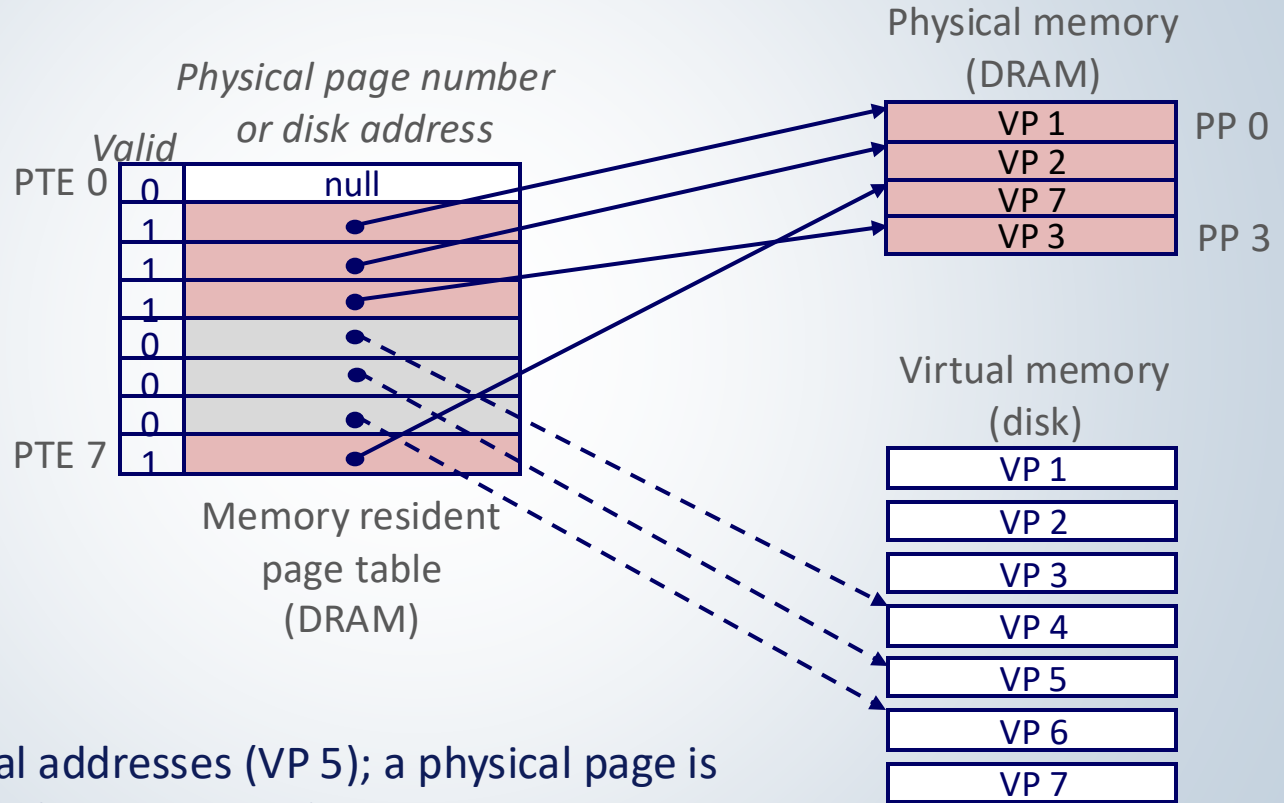


Handling Page Fault

- Page miss causes page fault (an exception)
- Page fault handler selects a victim to be evicted (here VP 4)
- Offending (faulting) instruction is restarted: page hit! On a **page fault**, the OS brings the page **from disk into DRAM** (or allocates a zero-filled page for anonymous memory), updates the PTE, and **restarts the faulting instruction**.



Allocating Pages





malloc() reserves virtual addresses (VP 5); a physical page is allocated on first access (demand-zero)

VM Address Translation

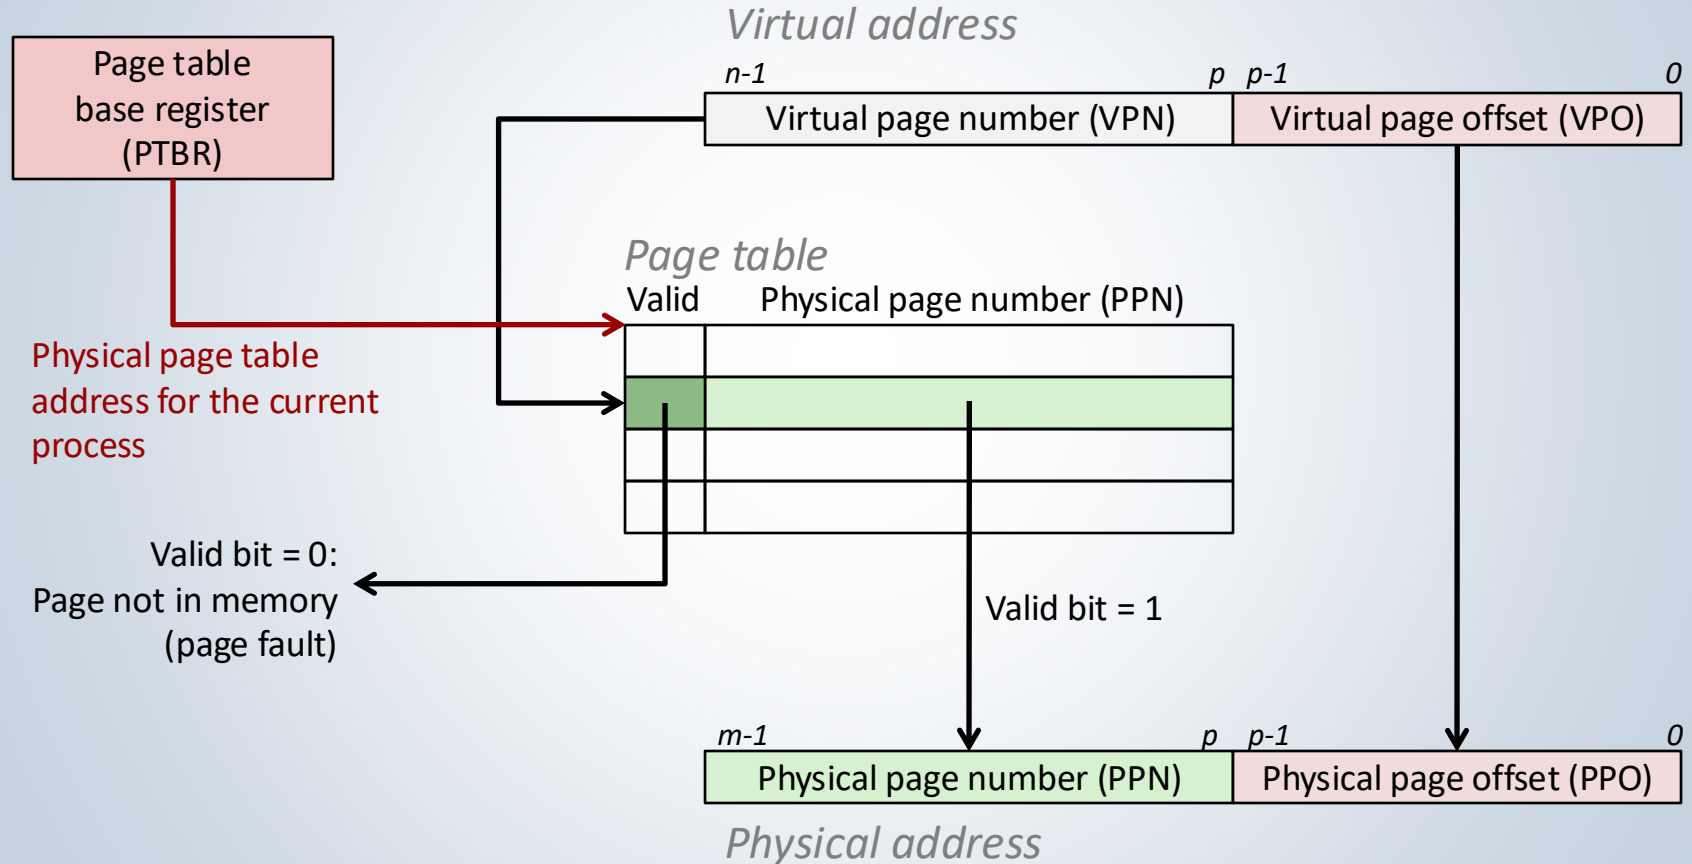
- Virtual Address Space
 - $V = \{0, 1, \dots, N-1\}$
- Physical Address Space
 - $P = \{0, 1, \dots, M-1\}$
- Address Translation
 - **$MAP: V \rightarrow P \cup \{\emptyset\}$**
 - For virtual address α :
 - **$MAP(\alpha) = \alpha'$** if data at virtual address α is at physical address α' in P
 - **$MAP(\alpha) = \emptyset$** if data at virtual address α is not in physical memory (page fault)

Breaking down virtual addresses

- Basic Parameters
 - $N = 2^n$: Number of addresses in virtual address space
 - $M = 2^m$: Number of addresses in physical address space. $m \leq n$ (usually much less)
 - $P = 2^p$: Page size (bytes)
- Components of the virtual address (VA)
 - Virtual page number (VPN): $n-p$ bits
 - Page Offset: p bits
- Components of the physical address (PA)
 - Physical page number (PPN): $m-p$ bits
 - Page Offset (same offset as VA): p bits



Address Translation with a Page Table



Virtual memory example

- Parameters

- Virtual addresses are 12-bits
- Physical addresses are 16-bits
- Page size is 64 bytes

For this toy example, which space is larger isn't important. **In real systems, virtual space is usually much larger than physical.**

1. How do we split Virtual Addresses into VPN and Offset?

11	10	9	8	7	6	5	4	3	2	1	0

Virtual memory example

- Parameters

- Virtual addresses are 12-bits
- Physical addresses are 16-bits
- Page size is 64 bytes

Mapping can be anything, which is bigger doesn't really matter!

- How do we split Virtual Addresses into VPN and Offset?

- Offset is based on page size: 64-bytes \Rightarrow 6 bits. All the rest are VPN

11	10	9	8	7	6	5	4	3	2	1	0
Virtual Page Number						Page Offset					

- How big are Physical Page Numbers?

Virtual memory example

- Parameters

- Virtual addresses are 12-bits
- Physical addresses are 16-bits
- Page size is 64 bytes

Mapping can be anything, which is bigger doesn't really matter!

- How do we split Virtual Addresses into VPN and Offset?

- Offset is based on page size: 64-bytes \Rightarrow 6 bits. All the rest are VPN

11	10	9	8	7	6	5	4	3	2	1	0
Virtual Page Number						Page Offset					

- How big are Physical Page Numbers? $16-6 = 10$ bits

Virtual memory example

- Parameters
 - Virtual addresses are 12-bits
 - Physical addresses are 16-bits
 - Page size is 64 bytes

11	10	9	8	7	6	5	4	3	2	1	0
Virtual Page Number						Page Offset					

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Physical Page Number										Page Offset					

- Translate:
- Virtual address: 0x3F0
 - Binary:
 - VPN:
 - Offset:

Virtual memory example

- Parameters
 - Virtual addresses are 12-bits
 - Physical addresses are 16-bits
 - Page size is 64 bytes

11	10	9	8	7	6	5	4	3	2	1	0
Virtual Page Number						Page Offset					

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Physical Page Number										Page Offset					

- Translate:
- Virtual address: 0x3F0
 - Binary: 0b001111110000
 - VPN: 0b001111
 - Offset: 0b110000

Virtual memory example

- Parameters

- Virtual addresses are 12-bits
- Physical addresses are 16-bits
- Page size is 64 bytes

11	10	9	8	7	6	5	4	3	2	1	0
Virtual Page Number						Page Offset					

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Physical Page Number										Page Offset					

- Translate:
- Virtual address: 0x3F0
 - Binary: 0b001111110000
 - VPN: 0b001111
 - Offset: 0b110000

VPN	PPN	Valid
0x00	0x123	1
0x01	0x156	1
0x02	0x143	1
0x03	0x16F	1
0x04	0x1FF	0
0x05	0x107	0
0x06	0x100	0
0x07	0x1C0	0
0x08	0x1D8	0
0x09	0x1BF	0
0x0A	0x000	1
0x0B	0x3FF	1
0x0C	0x308	0
0x0D	0x3FD	0
0x0E	0x111	1
0x0F	0x1F0	1

VPN	PPN	Valid
0x10	0x237	1
0x11	0x236	1
0x12	0x2B0	1
0x13	0x280	0
0x14	0x120	0
Continues on...		

- PPN:
- Offset:

Virtual memory example

- Parameters

- Virtual addresses are 12-bits
- Physical addresses are 16-bits
- Page size is 64 bytes

11	10	9	8	7	6	5	4	3	2	1	0
Virtual Page Number						Page Offset					

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Physical Page Number										Page Offset					

- Translate:
- Virtual address: 0x3F0
 - Binary: 0b001111110000
 - VPN: 0b001111
 - Offset: 0b110000

VPN	PPN	Valid
0x00	0x123	1
0x01	0x156	1
0x02	0x143	1
0x03	0x16F	1
0x04	0x1FF	0
0x05	0x107	0
0x06	0x100	0
0x07	0x1C0	0
0x08	0x1D8	0
0x09	0x1BF	0
0x0A	0x000	1
0x0B	0x3FF	1
0x0C	0x308	0
0x0D	0x3FD	0
0x0E	0x111	1
0x0F	0x1F0	1

VPN	PPN	Valid
0x10	0x237	1
0x11	0x236	1
0x12	0x2B0	1
0x13	0x280	0
0x14	0x120	0
Continues on...		

- PPN: 0b01 1111 0000
- Offset: 0b110000
- Physical address:

Virtual memory example

- Parameters

- Virtual addresses are 12-bits
- Physical addresses are 16-bits
- Page size is 64 bytes

11	10	9	8	7	6	5	4	3	2	1	0
Virtual Page Number						Page Offset					

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Physical Page Number										Page Offset					

- Translate:
- Virtual address: 0x3F0
 - Binary: 0b001111110000
 - VPN: 0b001111
 - Offset: 0b110000

VPN	PPN	Valid
0x00	0x123	1
0x01	0x156	1
0x02	0x143	1
0x03	0x16F	1
0x04	0x1FF	0
0x05	0x107	0
0x06	0x100	0
0x07	0x1C0	0
0x08	0x1D8	0
0x09	0x1BF	0
0x0A	0x000	1
0x0B	0x3FF	1
0x0C	0x308	0
0x0D	0x3FD	0
0x0E	0x111	1
0x0F	0x1F0	1

VPN	PPN	Valid
0x10	0x237	1
0x11	0x236	1
0x12	0x2B0	1
0x13	0x280	0
0x14	0x120	0
Continues on...		

- PPN: 0b0111110000
- Offset: 0b110000
- Physical address:
 - 0b0111110000110000
 - 0x7C30

- Parameters

- Virtual addresses are 12-bits
- Physical addresses are 16-bits
- Page size is 64 bytes

11	10	9	8	7	6	5	4	3	2	1	0
Virtual Page Number						Page Offset					

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Physical Page Number										Page Offset					

- Translate:
- Virtual address: 0x500
 - Binary:
 - VPN:
 - Offset:

VPN	PPN	Valid
0x00	0x123	1
0x01	0x156	1
0x02	0x143	1
0x03	0x16F	1
0x04	0x1FF	0
0x05	0x107	0
0x06	0x100	0
0x07	0x1C0	0
0x08	0x1D8	0
0x09	0x1BF	0
0x0A	0x000	1
0x0B	0x3FF	1
0x0C	0x308	0
0x0D	0x3FD	0
0x0E	0x111	1
0x0F	0x1F0	1

VPN	PPN	Valid
0x10	0x237	1
0x11	0x236	1
0x12	0x2B0	1
0x13	0x280	0
0x14	0x120	0
Continues on...		

- PPN:
- Offset:
- Physical address:

- Parameters

- Virtual addresses are 12-bits
- Physical addresses are 16-bits
- Page size is 64 bytes

11	10	9	8	7	6	5	4	3	2	1	0
Virtual Page Number						Page Offset					

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Physical Page Number										Page Offset					

- Translate:
- Virtual address: 0x500
 - Binary: 0b010100000000
 - VPN: 0b010100
 - Offset: 0b000000

VPN	PPN	Valid
0x00	0x123	1
0x01	0x156	1
0x02	0x143	1
0x03	0x16F	1
0x04	0x1FF	0
0x05	0x107	0
0x06	0x100	0
0x07	0x1C0	0
0x08	0x1D8	0
0x09	0x1BF	0
0x0A	0x000	1
0x0B	0x3FF	1
0x0C	0x308	0
0x0D	0x3FD	0
0x0E	0x111	1
0x0F	0x1F0	1

VPN	PPN	Valid
0x10	0x237	1
0x11	0x236	1
0x12	0x2B0	1
0x13	0x280	0
0x14	0x120	0
Continues on...		

- PPN: INVALID
- Offset:
- Physical address:
 - Page Fault

- Parameters

- Virtual addresses are 12-bits
- Physical addresses are 16-bits
- Page size is 64 bytes

11	10	9	8	7	6	5	4	3	2	1	0
Virtual Page Number						Page Offset					

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Physical Page Number										Page Offset					

- Translate:
- Virtual address: 0x0D6
 - Binary:
 - VPN:
 - Offset:

VPN	PPN	Valid
0x00	0x123	1
0x01	0x156	1
0x02	0x143	1
0x03	0x16F	1
0x04	0x1FF	0
0x05	0x107	0
0x06	0x100	0
0x07	0x1C0	0
0x08	0x1D8	0
0x09	0x1BF	0
0x0A	0x000	1
0x0B	0x3FF	1
0x0C	0x308	0
0x0D	0x3FD	0
0x0E	0x111	1
0x0F	0x1F0	1

VPN	PPN	Valid
0x10	0x237	1
0x11	0x236	1
0x12	0x2B0	1
0x13	0x280	0
0x14	0x120	0
Continues on...		

- PPN:
- Offset:
- Physical address:

- Parameters

- Virtual addresses are 12-bits
- Physical addresses are 16-bits
- Page size is 64 bytes

11	10	9	8	7	6	5	4	3	2	1	0
Virtual Page Number						Page Offset					

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
Physical Page Number										Page Offset					

- Translate:
- Virtual address: 0x0D6
 - Binary: 0b000011010110
 - VPN: 0b000011
 - Offset: 0b010110


VPN	PPN	Valid
0x00	0x123	1
0x01	0x156	1
0x02	0x143	1
0x03	0x16F	1
0x04	0x1FF	0
0x05	0x107	0
0x06	0x100	0
0x07	0x1C0	0
0x08	0x1D8	0
0x09	0x1BF	0
0x0A	0x000	1
0x0B	0x3FF	1
0x0C	0x308	0
0x0D	0x3FD	0
0x0E	0x111	1
0x0F	0x1F0	1

VPN	PPN	Valid
0x10	0x237	1
0x11	0x236	1
0x12	0x2B0	1
0x13	0x280	0
0x14	0x120	0
Continues on...		

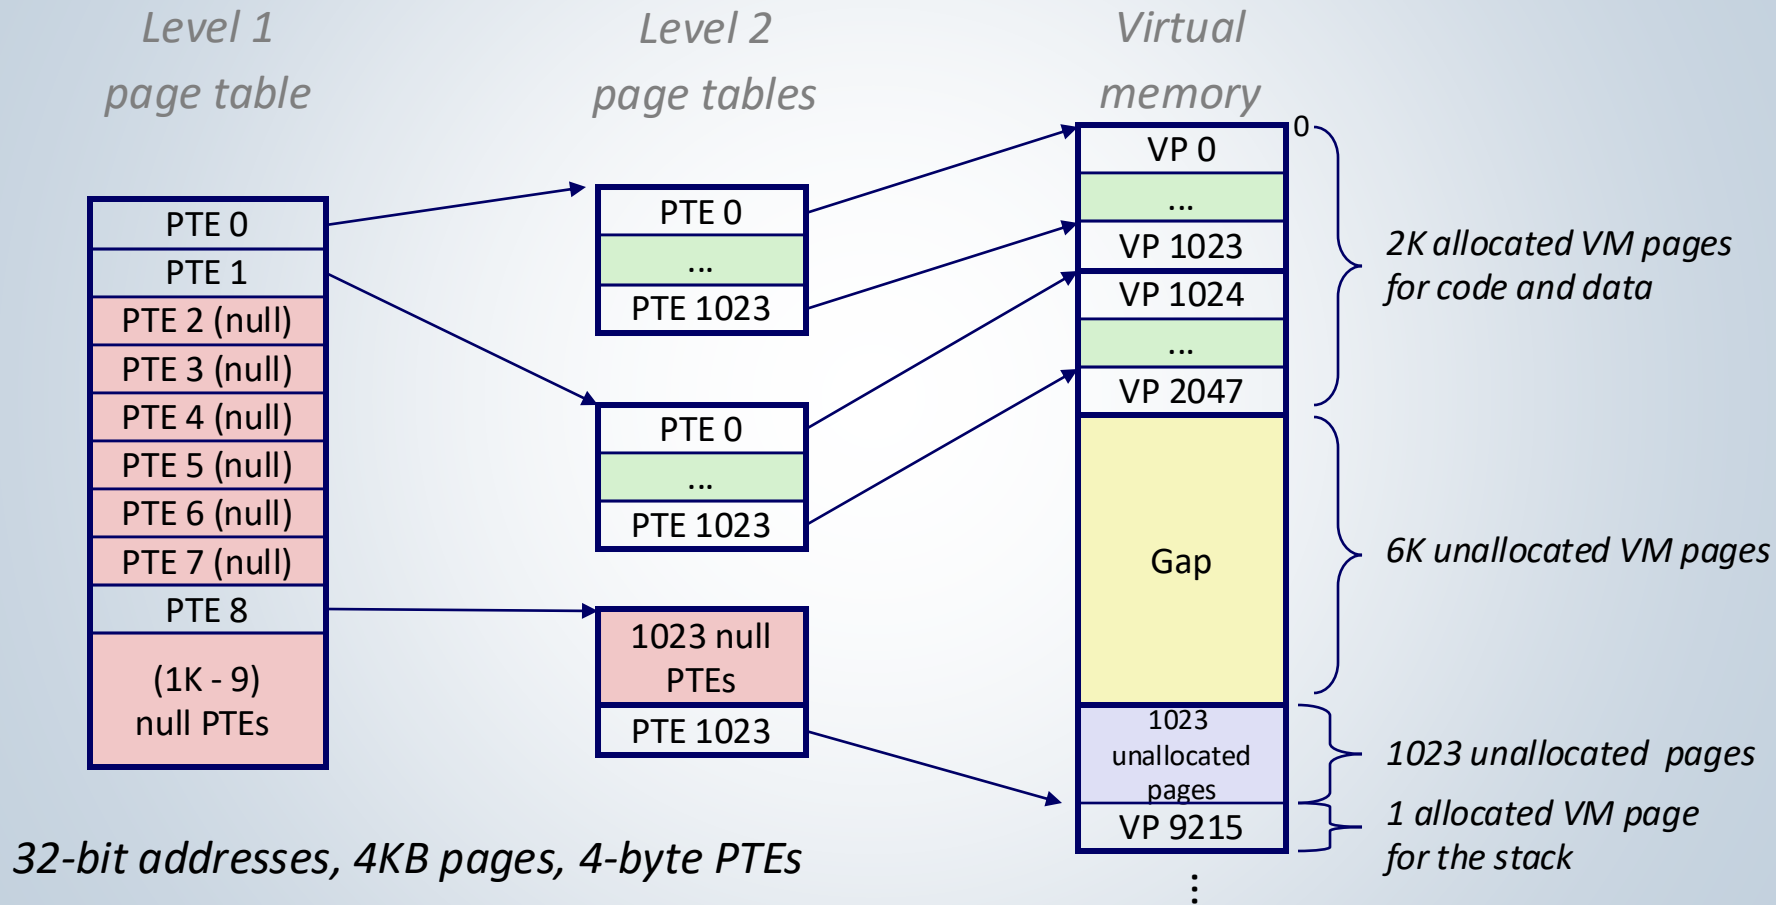
- PPN: 0b010 110 1111
- Offset: 0b010110
- Physical address:
 - 0b0101101111010110
 - 0x5BD6

Multi-Level Page Tables

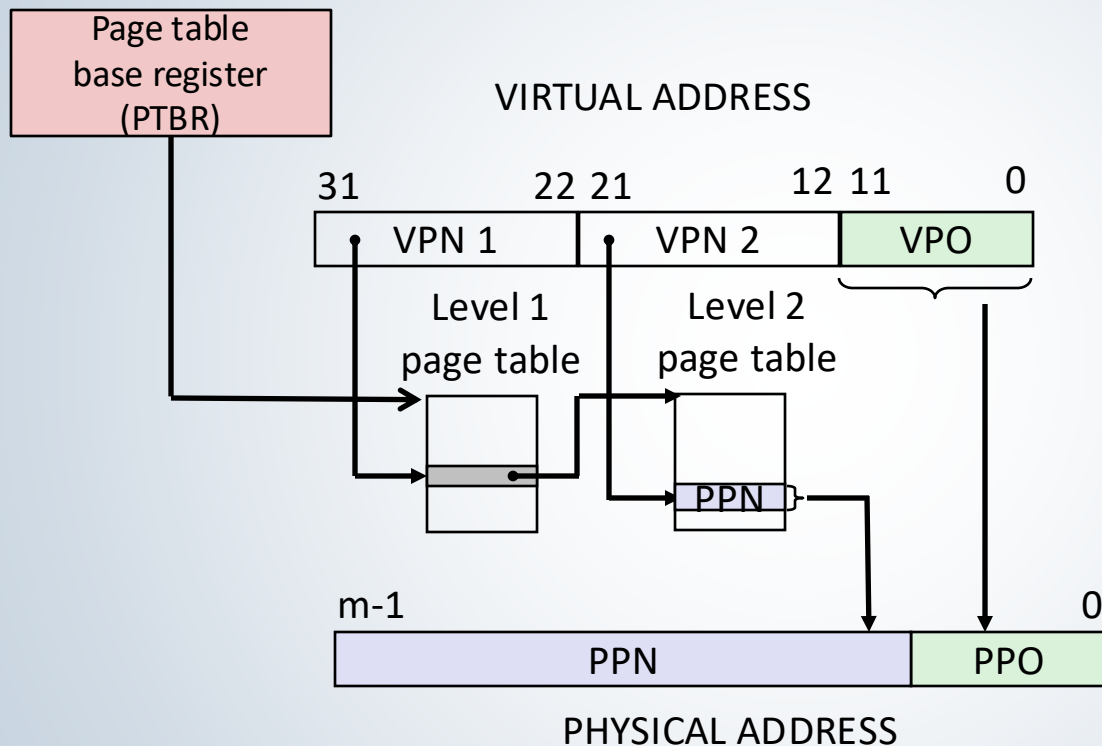
- Suppose:
 - 4KB (2^{12}) page size, 48-bit address space, 8-byte PTE (2^3)
- Problem:
 - Would need a 512 GB page table!
 - $2^{48} * 2^{-12} * 2^3 = 2^{39}$ bytes
- Common solution: Multi-level page table
- Example: 2-level page table
 - Level 1 table: each PTE points to a page table (always memory resident)
 - Level 2 table: each PTE points to a data page (paged in and out like any other data)
 - Huge advantage: Allocate an L2 table when some virtual pages in that region are actually used



Two-Level Page-Table Hierarchy



Translating with a 2-Level Page Table



Multi-level page table (example)

- Consider a 3-level page table. A page table has 64 entries. If the page size is 4K bytes, how much is the size of the virtual address space?

Multi-level page table (example)

- Consider a 3-level page table. A page table has 64 entries. If the page size is 4K bytes, how much is the size of the virtual address space?
- 64 entries per table: 6 bits of index per level
- Page size = 4KB, 2^{12} \rightarrow 12 offset bits
- Bits view
 - 3 levels * 6 bits = 18 + offset = 30 bits
 - The virtual address space size is 2^{30} bytes = **1 GiB**

Multi-level page table (example)

- Another view:
- L3 table maps 64 pages: $64 \times 4\text{KB} = 256\text{KB}$
- L2 table has 64 entries, each mapping to an L3: $64 \times 256\text{KB} = 16\text{MB}$
- L1 has 64 entries, each mapping to an L2: $64 \times 16\text{MB} = 1\text{GB}$