



CS3281 / CS5281

# Concurrent Programming

Will Hedgecock  
Sandeep Neema  
Bryan Ward

*\*Some lecture slides borrowed and adapted from CMU's  
"Computer Systems: A Programmer's Perspective"*



Tel (615) 343-7472 | Fax (615) 343-7440  
1025 16th Avenue South Nashville, TN 37212  
[www.isis.vanderbilt.edu](http://www.isis.vanderbilt.edu)



# Motivation

- Idea: Use multiple processes to “collaborate” to accomplish work faster
  - Divide up work between two or more threads
    - Single-core hardware performance advances slowly, but we are seeing more cores and parallel architectures
  - Run concurrently on multiple cores
    - Fun fact: it’s actually possible to get “super-linear” speedups!
  - A process may be able to run while another waits for some event
    - e.g., waiting on a socket for a network packet (as we will see later)
- What challenges might we face?

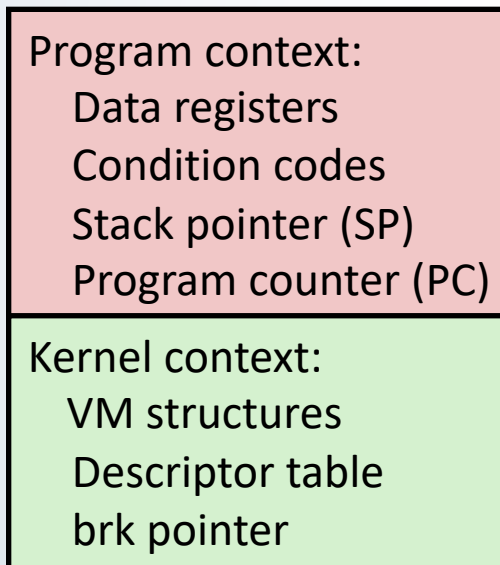
# Terminology

- Concurrent programming: a programming paradigm in which multiple tasks are executed in overlapping periods
  - Enables multiple operations to be executed out-of-order without adversely affecting final outcome
- Benefits
  - Speed & responsiveness
  - Better resource utilization
  - Scalability in modern architectures

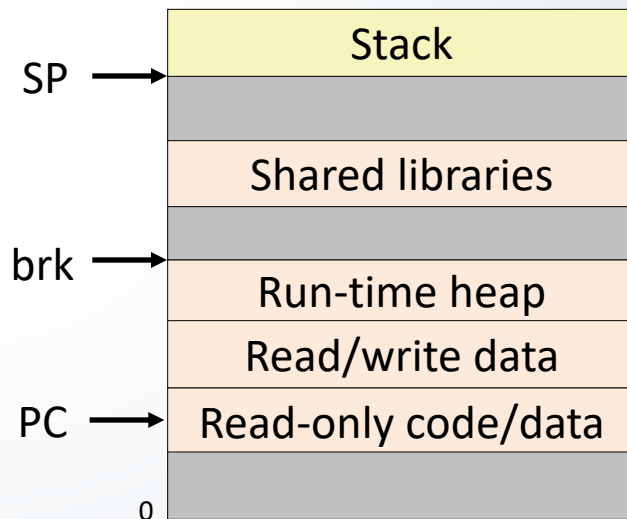
# Traditional View of a Process

- Process = process context + code, data, and stack

## Process context



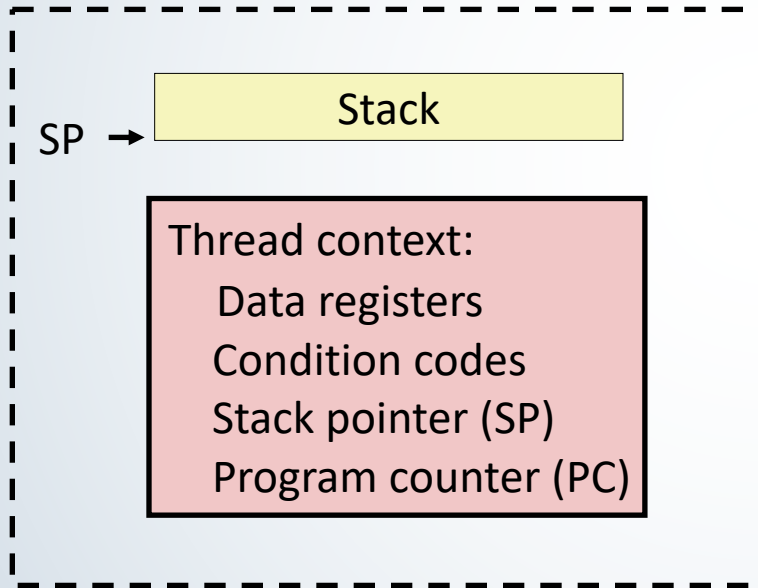
## Code, data, and stack



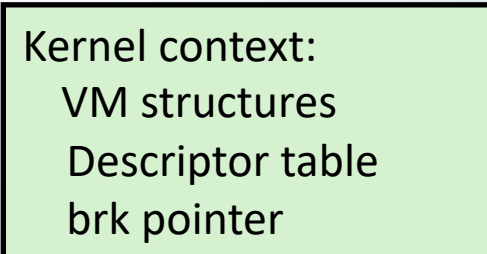
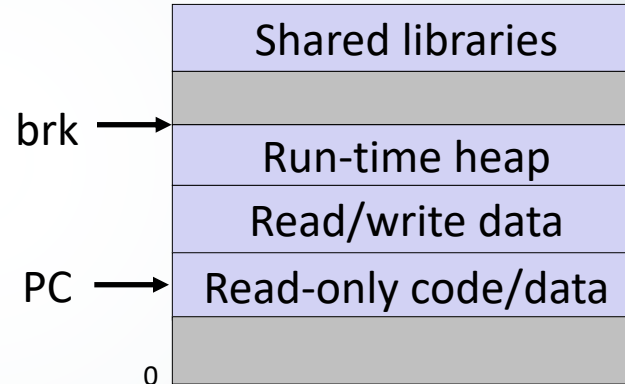
# Alternate View of a Process

- Process = thread + code, data, and kernel context

Thread (main thread)



Code, data, and kernel context



# A Process with Multiple Threads

- Multiple threads can be associated with a process
  - Each thread has its own logical control flow
  - Each thread shares the same code, data, and kernel context (unlike processes)
  - Each thread has its own stack for local variables
    - but not protected from other threads – why?
  - Each thread has its own thread id (TID)

## Thread 1 (main thread)

stack 1

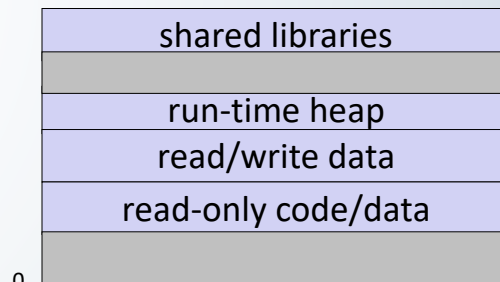
Thread 1 context:  
Data registers  
Condition codes  
SP1  
PC1

## Thread 2 (peer thread)

stack 2

Thread 2 context:  
Data registers  
Condition codes  
SP2  
PC2


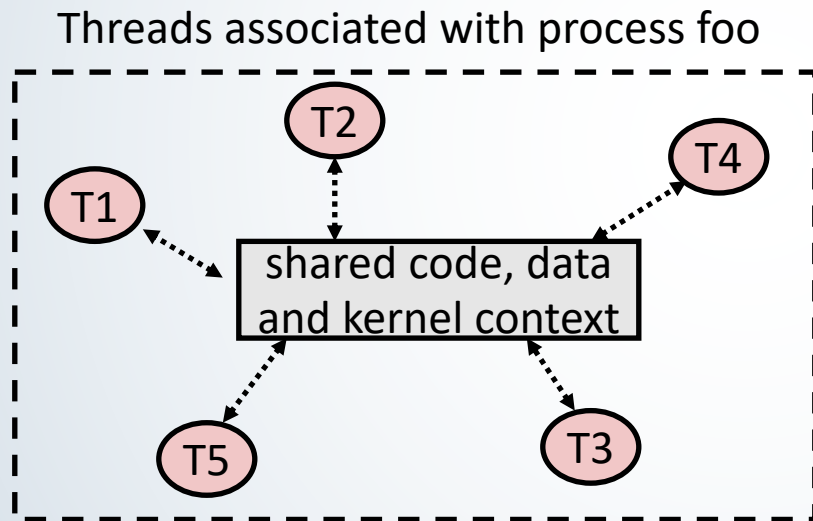
## Shared code and data



Kernel context:  
VM structures  
Descriptor table  
brk pointer

# Logical View of Threads

- Threads associated with process form a pool of peers
  - Unlike processes, which form a tree hierarchy




# Concurrent Threads

- Two threads are *concurrent* if their flows overlap in time
- Otherwise, they are sequential

- Examples:


- Concurrent: A & B, A&C
- Sequential: B & C






# Concurrent Thread Execution

- Single Core Processor
  - Simulate parallelism by time slicing



- Multi-Core Processor
  - Can have true parallelism



# Threads vs. Processes

- How threads and processes are similar
  - Each has its own logical control flow
  - Each can run concurrently with others (possibly on different cores)
  - Each is context switched
- How threads and processes are different
  - Threads share all code and data (except local stacks)
    - Processes (typically) do not
  - Threads are somewhat less expensive than processes
    - Process control (creating and reaping) twice as expensive as thread control
    - Linux numbers:
      - ~20K cycles to create and reap a process
      - ~10K cycles (or less) to create and reap a thread
    - Switching between threads of the same process is faster than switching between processes

# Posix Threads (pthreads) Interface (not xv6)

- *Pthreads*: Standard interface for ~60 functions that manipulate threads from C programs
  - Creating and reaping threads
    - `pthread_create()`
    - `pthread_join()`
  - Determining your thread ID
    - `pthread_self()`
  - Terminating threads
    - `pthread_cancel()`
    - `pthread_exit()`
    - `exit()` [terminates all threads] , `RET` [terminates current thread]
  - Synchronizing access to shared variables
    - `pthread_mutex_init`
    - `pthread_mutex_[un]lock`

# The pthreads "Hello, World" Program

```
// hello.c – Pthreads "hello, world" program
```

```
#include "csapp.h"
```

```
void *thread(void *vargp);
```

```
int main()
```

```
{  
    pthread_t tid;  
    Pthread_create(&tid, NULL, thread, NULL);  
    Pthread_join(tid, NULL);  
    exit(0);  
}
```

hello.c

Thread ID

Thread attributes  
(usually NULL)

Thread routine

Thread arguments  
(void \*p)

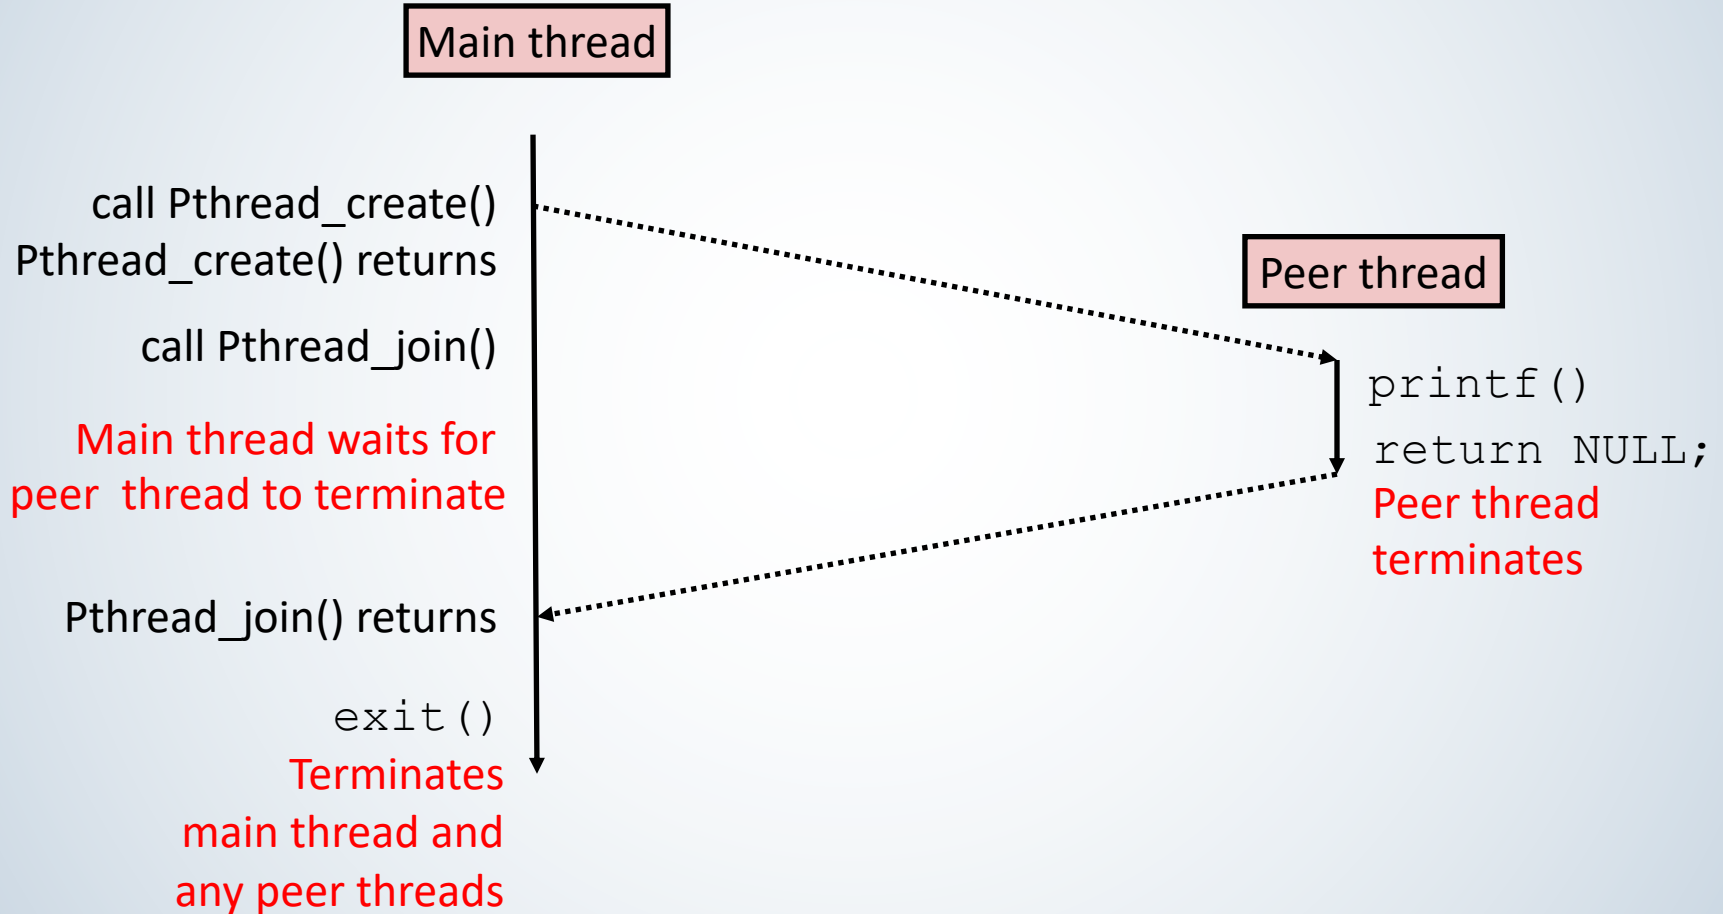
Return value  
(void \*\*p)

```
void *thread(void *vargp) /* thread routine */
```

```
{  
    printf("Hello, world!\n");  
    return NULL;  
}
```

hello.c

# Execution of Threaded “Hello, World”



# Issues with Threads

- Must run “detached” to avoid memory leak
  - At any point in time, a thread is either *joinable* or *detached*
  - *Joinable* thread can be reaped and killed by other threads
    - must be reaped (with `pthread_join`) to free memory resources
  - *Detached* thread cannot be reaped or killed by other threads
    - resources are automatically reaped on termination
  - Default state is joinable
    - use `pthread_detach(pthread_self())` to make detached
- Must be careful to avoid unintended sharing
  - For example, passing pointer to main thread’s stack
    - `pthread_create(&tid, NULL, thread, (void *)buffer);`
- All functions called by concurrent thread must be *thread-safe*
  - (next lecture)

# Pros and Cons of Thread-Based Designs

- + Easy to share data structures between threads
  - e.g., logging information, file cache
- + Threads are more efficient than processes
- – Unintentional sharing can introduce subtle and hard-to-reproduce errors!
  - The ease with which data can be shared is both the greatest strength and the greatest weakness of threads
  - Hard to know which data should be shared vs. private
  - Hard to detect by testing
    - Probability of bad “race” outcome very low
    - But nonzero!
  - Future lectures