



CS3281 / CS5281

# Concurrency Bugs

Will Hedgecock  
Sandeep Neema  
Bryan Ward

*\*Some lecture slides borrowed and adapted from “Operating Systems: 3 Easy Pieces”  
and CMU’s “Computer Systems: A Programmer’s Perspective”*



Tel (615) 343-7472 | Fax (615) 343-7440  
1025 16th Avenue South Nashville, TN 37212  
[www.isis.vanderbilt.edu](http://www.isis.vanderbilt.edu)



# Non-Deadlock Bugs: Atomicity Violation

- Thread 1 is interrupted before it runs fputs() function
- Dereference the null pointer exception inside fputs()
- Memory access (thd->proc\_info) needs to be protected

```
Thread1::  
if (thd->proc_info) {  
    ...  
    fputs(thd->proc_info , ...);  
    ...  
}  
  
Thread2::  
thd->proc_info = NULL;
```

# Non-Deadlock Bugs: Atomicity Violation

- Fix: place `thd->proc_info` within lock and unlock routines

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;

Thread1::
pthread_mutex_lock(&lock);
if (thd->proc_info) {
    ...
    fputs(thd->proc_info , ...);
    ...
}
pthread_mutex_unlock(&lock);

Thread2::
pthread_mutex_lock(&lock);
thd->proc_info = NULL;
pthread_mutex_unlock(&lock);
```

# Non-Deadlock Bugs: Order Violation

- Thread 2 runs before Thread 1
  - mThread is not initialized
  - Null-pointer dereference

```
Thread1::  
void init() {  
    mThread = PR_CreateThread(mMain, ...);  
}  
  
Thread2::  
void mMain(...) {  
    mState = mThread->State  
}
```

# Non-Deadlock Bugs: Order Violation

- Fix: Use condition variables
- Enforce the order of execution between memory accesses

```
pthread_mutex_t mtLock = PTHREAD_MUTEX_INITIALIZER;  
pthread_cond_t mtCond = PTHREAD_COND_INITIALIZER;  
int mtInit = 0;
```

**Thread 1::**

```
void init(){  
    ...  
    mThread = PR_CreateThread(mMain,...);  
  
    // signal that the thread has been created.  
    pthread_mutex_lock(&mtLock);  
    mtInit = 1;  
    pthread_cond_signal(&mtCond);  
    pthread_mutex_unlock(&mtLock);  
    ...  
}
```

**Thread2::**

```
void mMain(...){  
    ...  
}
```


```
// wait for the thread to be initialized ...  
pthread_mutex_lock(&mtLock);  
while(mtInit == 0)  
    pthread_cond_wait(&mtCond, &mtLock);  
pthread_mutex_unlock(&mtLock);  
  
mState = mThread->State;  
...  
}
```

# Deadlock

- Def: A process is *deadlocked* iff it is waiting for a condition that will never be true
- Typical Scenario
  - Processes 1 and 2 needs two resources (A and B) to proceed
  - Process 1 acquires A, waits for B
  - Process 2 acquires B, waits for A
  - Both will wait forever!

# Deadlock Bugs

- Deadlock: A thread holds a lock and waits for another lock.



*Thread 1:*  
`pthread_mutex_lock (L1);`  
context switch to Thread 2  
`pthread_mutex_lock (L2);`

*Thread 2;*  
`pthread_mutex_lock (L2);`  
`pthread_mutex_lock (L1);`

# Deadlock Bugs

- Encapsulation does not work well with locking
- Vector class in Java
- *Thread 1:*            *Thread 2:*
- Vector v1, v2;
- v1.addAll (v2);    v2.addAll (v1);
- *addAll()* needs to be thread safe.
- Thread 1 acquires a lock for v1.
- Thread 2 acquires a lock for v2 at the same time.



# Deadlock Bugs

- Conditions for deadlocks
  - Mutual exclusion: Thread grabs a lock
  - Hold-and-wait: Thread holds the lock and waits to acquire an additional lock
  - No preemption: Lock that is held cannot be taken away from the thread
  - Circular wait: *hold-and-wait* for a circular chain of threads
- Deadlock does not occur if **any** of the above conditions is not met

# Deadlock Bugs

- Prevention for: Circular Wait
  - Total ordering (two locks): acquire L1 before L2
  - Partial ordering (multiple locks): group lock acquisition ordering
    - E.g., acquire L1 before L2, acquire L3 before L4
- Prevention for: Hold-and-Wait
  - Acquire all locks at once
    - pthread\_mutex\_lock (prevention);
    - pthread\_mutex\_lock (L1);
    - pthread\_mutex\_lock (L2);
    - ...
    - pthread\_mutex\_unlock (prevention);
  - Disadvantage: must know all needed locks ahead of time

# Deadlock Bugs

- Prevention for: No Preemption
- Use `pthread_mutex_trylock()`
  - Acquire lock if it is available or return an error code (lock is already held)
  - Preemption: release the ownership of a lock

```
top:
    lock(L1);
    if( tryLock(L2) == -1 ){
        unlock(L1);
        goto top;
    }
```

- Livelock
  - Another thread holds L2 and attempts to acquire L1.
  - Both threads execute their code blocks at the same time.
- Solution: adding a random delay before jumping back

# Deadlock Bugs

- Problems with trylock
  - Memory allocated after acquiring L1 should be released when acquiring L2 fails.
  - trylock does not preempt the ownership of a lock. It allows a thread to give back the ownership voluntarily.

# Deadlock Bugs

- Prevention for: Mutual Exclusion
  - Lock-free data structures
- Create atomic functions based on hardware instructions

```
int CompareAndSwap(int *address, int expected, int new){
    if(*address == expected){
        *address = new;
        return 1; // success
    }
    return 0;
}

void AtomicIncrement(int *value, int amount){
    do{
        int old = *value;
    }while( CompareAndSwap(value, old, old+amount)==0);
}
```

# Deadlock Bugs

- List insertion
- Race condition occurs if called by multiple threads

```
void insert(int value){
    node_t * n = malloc(sizeof(node_t));
    assert( n != NULL );
    n->value = value ;
    n->next = head;
    head    = n; ← interrupt
}


void insert(int value) {
    node_t *n = malloc(sizeof(node_t));
    assert(n != NULL);
    n->value = value;
    do {
        n->next = head;
    } while (CompareAndSwap(&head, n->next, n));
}
```

# Deadlock Bugs

- Avoid deadlock
- Find out what locks are acquired by what threads

	T1	T2	T3	T4
L1	yes	yes	no	no
L2	yes	yes	yes	no

- Scheduler does not run T1 and T2 at the same time.
- T3 grabs only one lock. It can run with T1.



# Deadlock Bugs

	T1	T2	T3	T4
L1	yes	yes	yes	no
L2	yes	yes	yes	no



- Decrease concurrency
- Trade-off between performance and deadlock avoidance
- It's impractical to gain a priori knowledge of lock acquisition for most applications.



# Deadlock Bugs

- Allow deadlocks to occur
  - Find mechanisms to detect and recover
  - For example, a recovery mechanism reboots a system after detecting a cycle

# Deadlocking with Semaphores

```
int main()
{
    pthread_t tid[2];
    Sem_init(&mutex[0], 0, 1); /*
mutex[0] = 1 */
    Sem_init(&mutex[1], 0, 1); /*
mutex[1] = 1 */
    Pthread_create(&tid[0], NULL, count,
(void*) 0);
    Pthread_create(&tid[1], NULL, count,
(void*) 1);
    Pthread_join(tid[0], NULL);
    Pthread_join(tid[1], NULL);
    printf("cnt=%d\n", cnt);
}
```

```
void *count(void *vargp)
{
    int i;
    int id = (int) vargp;
    for (i = 0; i < NITERS; i++) {
        P(&mutex[id]); P(&mutex[1-id]);
        cnt++;
        V(&mutex[id]); V(&mutex[1-id]);
    }
    return NULL;
}
```

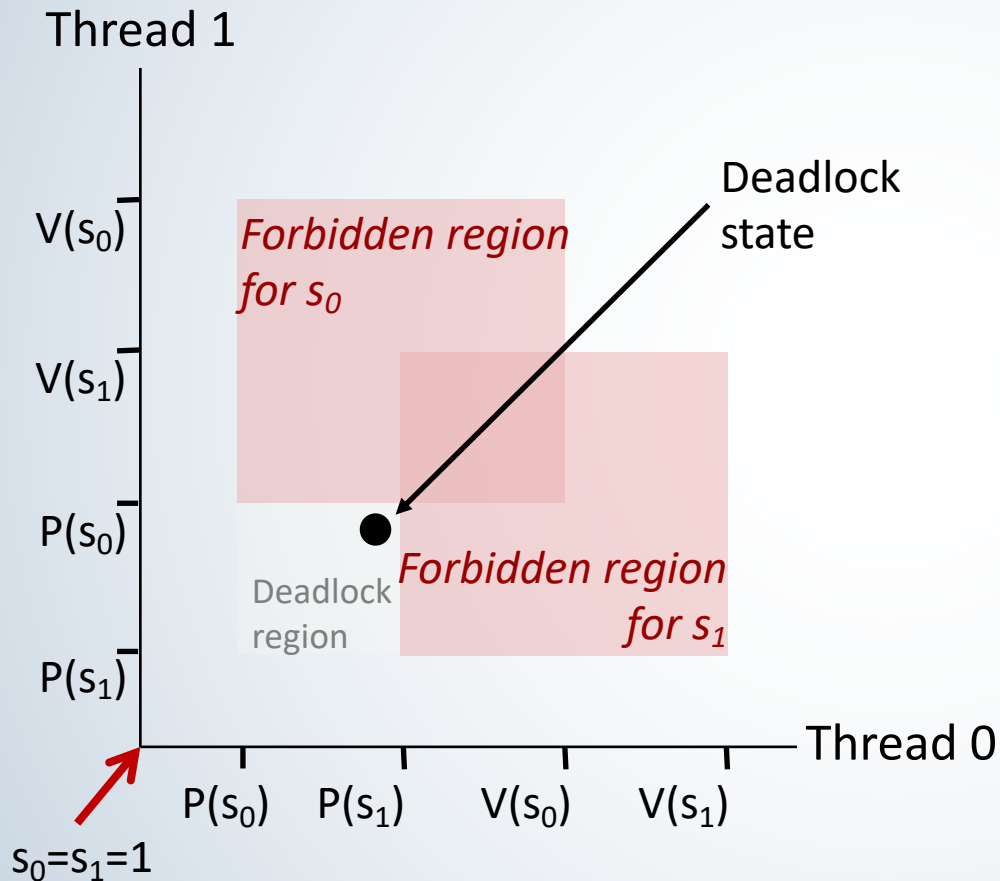
Tid[0]:

P(s<sub>0</sub>);  
P(s<sub>1</sub>);  
cnt++;  
V(s<sub>0</sub>);  
V(s<sub>1</sub>);

Tid[1]:

P(s<sub>1</sub>);  
P(s<sub>0</sub>);  
cnt++;  
V(s<sub>1</sub>);  
V(s<sub>0</sub>);

# Deadlock Visualization in a Progress Graph



Locking introduces the potential for *deadlock*: waiting for a condition that will never be true

Any trajectory that enters the *deadlock region* will eventually reach the *deadlock state*, waiting for either  $s_0$  or  $s_1$  to become nonzero

Other trajectories luck out and skirt the deadlock region

Unfortunate fact: deadlock is often nondeterministic (race)

# Avoiding Deadlock


```
int main()
{
    pthread_t tid[2];
    Sem_init(&mutex[0], 0, 1);  /* mutex[0] = 1 */
    Sem_init(&mutex[1], 0, 1);  /* mutex[1] = 1 */
    Pthread_create(&tid[0], NULL, count, (void*) 0);
    Pthread_create(&tid[1], NULL, count, (void*) 1);
    Pthread_join(tid[0], NULL);
    Pthread_join(tid[1], NULL);
    printf("cnt=%d\n", cnt);
    exit(0);
}
```

```
void *count(void *vargp)
{
    int i;
    int id = (int) vargp;
    for (i = 0; i < NITERS; i++) {
        P(&mutex[0]); P(&mutex[1]);
        cnt++;
        V(&mutex[id]); V(&mutex[1-id]);
    }
    return NULL;
}
```

Tid[0]:  
P(s0);  
P(s1);  
cnt++;  
V(s0);  
V(s1);

Tid[1]:  
P(s0);  
P(s1);  
cnt++;  
V(s1);  
V(s0);

# Avoiding Deadlock in a Progress Graph



No way for trajectory to get stuck

Processes acquire locks in same order

Order in which locks released immaterial

# Summary

- Common solutions for deadlocks
  - Prevent deadlock
    - Order lock acquisition
    - Acquire lock atomically
    - Release lock voluntarily
    - Build lock-free atomic operations
  - Avoid deadlock
    - Schedule threads with global knowledge of lock acquisitions