

Lecture 11: Concurrency Bugs

CS 3281

Daniel Balasubramanian, Shervin Hajiamini

Non-Deadlock Bugs

- Atomicity violation
- Thread 1 is interrupted before it runs fputs() function
- Dereference the null pointer exception inside fputs()
- Memory access (thd->proc_info) needs to be serialized.

Thread1 ::

```
if (thd->proc_info) {  
    ...  
    fputs(thd->proc_info , ...);  
    ...  
}
```

Thread2 ::

```
thd->proc_info = NULL;
```

Non-Deadlock Bugs

Fix: place `thd->proc_info` within lock and unlock routines

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

Thread1 ::

```
pthread_mutex_lock(&lock);  
if (thd->proc_info) {  
    ...  
    fputs(thd->proc_info , ...);  
    ...  
}  
pthread_mutex_unlock(&lock);
```

Thread2 ::

```
pthread_mutex_lock(&lock);  
thd->proc_info = NULL;  
pthread_mutex_unlock(&lock);
```

Non-Deadlock Bugs

- Order violation
- Thread 2 runs after creation.
 - mThread is not initialized.
 - Null-pointer dereference

Thread1::

```
void init() {  
    mThread = PR_CreateThread(mMain, ...);  
}
```

Thread2::

```
void mMain(...) {  
    mState = mThread->State  
}
```

Non-Deadlock Bugs

- Fix: Use condition variables
- Enforce the order of execution between memory accesses

```
pthread_mutex_t mtLock = PTHREAD_MUTEX_INITIALIZER;  
pthread_cond_t mtCond = PTHREAD_COND_INITIALIZER;  
int mtInit = 0;
```

Thread 1::

```
void init() {  
    ...  
    mThread = PR_CreateThread(mMain, ...);  
  
    // signal that the thread has been created.  
    pthread_mutex_lock(&mtLock);  
    mtInit = 1;  
    pthread_cond_signal(&mtCond);  
    pthread_mutex_unlock(&mtLock);  
    ...  
}
```

Thread2::

```
void mMain(...) {  
    ...
```

Non-Deadlock Bugs

```
// wait for the thread to be initialized ...
pthread_mutex_lock(&mtLock);
while(mtInit == 0)
    pthread_cond_wait(&mtCond, &mtLock);
pthread_mutex_unlock(&mtLock);

mState = mThread->State;
...
}
```

Another worry: Deadlock

- Def: A process is *deadlocked* iff it is waiting for a condition that will never be true

- Typical Scenario
 - Processes 1 and 2 needs two resources (A and B) to proceed
 - Process 1 acquires A, waits for B
 - Process 2 acquires B, waits for A
 - Both will wait forever!

Deadlock Bugs

- Deadlock

- A thread holds a lock and waits for another lock.

Thread 1:

`pthread_mutex_lock(L1);`


context switch to Thread 2

`pthread_mutex_lock(L2);`

Thread 2;

`pthread_mutex_lock(L2);`

`pthread_mutex_lock(L1);`



Deadlock Bugs

- Encapsulation does not work well with locking.
- Vector class in Java

Thread 1: *Thread 2:*
Vector v1, v2;
v1.addAll (v2); v2.addAll (v1);

- *addAll ()* needs to be thread safe.
Thread 1 acquires a lock for v1.
Thread 2 acquires a lock for v2 at the same time.

Deadlock Bugs

- Conditions for deadlocks
 - Mutual exclusion: a thread grabs a lock.
 - Hold-and-wait: a thread holds a lock and waits for acquiring an additional lock.
 - No preemption: a lock that is held by a thread cannot be taken away from the thread.
 - Circular wait: *hold-and-wait* for a circular chain of threads.
- Deadlock does not occur if any of the conditions is not met.

Deadlock Bugs

- Prevent circular wait
 - Total ordering (two locks): acquire L1 before L2.
 - Partial ordering (multiple locks): group lock acquisition ordering
 - acquire L1 before L2, acquire L3 before L4
- Prevent Hold-and-Wait
 - Acquire all locks at once

```
pthread_mutex_lock (prevention);  
pthread_mutex_lock (L1);  
pthread_mutex_lock (L2);  
...  
pthread_mutex_unlock (prevention);
```

➤ Disadvantage: must know ahead of time all needed locks

Deadlock Bugs

- Prevent no preemption
- Use `pthread_mutex_trylock()`
 - Hold lock if it is available or return an error code (lock is already held)
 - Preemption: returning the ownership of a lock

```
top:
    lock(L1);
    if( tryLock(L2) == -1 ) {
        unlock(L1);
        goto top;
    }
```

- Livelock
 - Another thread holds L2 and attempts to acquire L1.
 - Both threads execute their code blocks at the same time.
- Solution: adding a random delay before jumping back

Deadlock Bugs

- Problems with trylock
 - Memory allocated after acquiring L1 should be released when acquiring L2 fails.
 - trylock does not preempt the ownership of a lock. It allows a thread to give back the ownership voluntarily.

Deadlock Bugs

- Prevent mutual exclusion
 - Lock-free data structures
- Create atomic functions based on hardware instructions

```
int CompareAndSwap(int *address, int expected, int new) {  
    if(*address == expected) {  
        *address = new;  
        return 1; // success  
    }  
    return 0;  
}
```

```
void AtomicIncrement(int *value, int amount) {  
    do {  
        int old = *value;  
    } while ( CompareAndSwap(value, old, old+amount) == 0 );  
}
```

Deadlock Bugs

- List insertion
- Race condition occurs if called by multiple threads

```
void insert(int value) {  
    node_t * n = malloc(sizeof(node_t));  
    assert( n != NULL );  
    n->value = value ;  
    n->next  = head;  
    head     = n;    ← interrupt  
}
```


```
void insert(int value) {  
    node_t *n = malloc(sizeof(node_t));  
    assert(n != NULL);  
    n->value = value;  
    do {  
        n->next = head;  
    } while (CompareAndSwap(&head, n->next, n));  
}
```

Deadlock Bugs

- Avoid deadlock
- Find out what locks are acquired by what threads


	T1	T2	T3	T4
L1	yes	Yes	no	no
L2	yes	yes	yes	no

- Scheduler does not run T1 and T2 at the same time.
- T3 grabs only one lock. It can run with T1.



Deadlock Bugs

	T1	T2	T3	T4
L1	yes	Yes	yes	no
L2	yes	yes	yes	no



- Decrease concurrency
- Trade-off between performance and deadlock avoidance
- It's impractical to gain a priori knowledge of lock acquisition for most applications.

Deadlock Bugs

- Allow deadlocks to occur
 - Find mechanisms to detect and recover
 - For example, a recovery mechanism reboots a system after detecting a cycle

Deadlocking With Semaphores

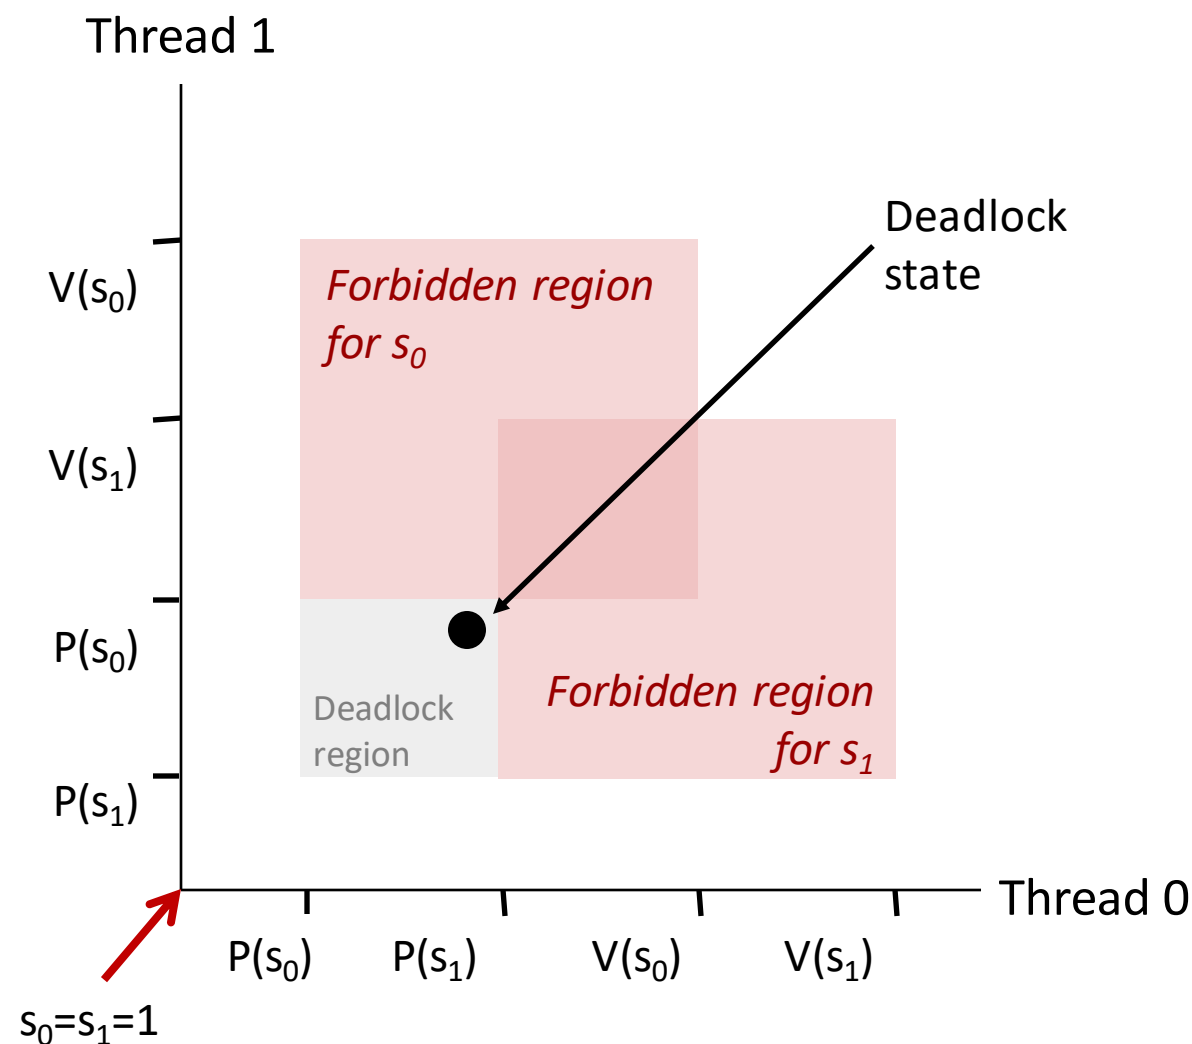
```
int main()
{
    pthread_t tid[2];
    Sem_init(&mutex[0], 0, 1); /* mutex[0] = 1 */
    Sem_init(&mutex[1], 0, 1); /* mutex[1] = 1 */
    Pthread_create(&tid[0], NULL, count, (void*) 0);
    Pthread_create(&tid[1], NULL, count, (void*) 1);
    Pthread_join(tid[0], NULL);
    Pthread_join(tid[1], NULL);
    printf("cnt=%d\n", cnt);
    exit(0);
}
```

```
void *count(void *vargp)
{
    int i;
    int id = (int) vargp;
    for (i = 0; i < NITERS; i++) {
        P(&mutex[id]); P(&mutex[1-id]);
        cnt++;
        V(&mutex[id]); V(&mutex[1-id]);
    }
    return NULL;
}
```

Tid[0]:
P(s₀);
P(s₁);
cnt++;
V(s₀);
V(s₁);

Tid[1]:
P(s₁);
P(s₀);
cnt++;
V(s₁);
V(s₀);

Deadlock Visualized in Progress Graph



Locking introduces the potential for *deadlock*: waiting for a condition that will never be true

Any trajectory that enters the *deadlock region* will eventually reach the *deadlock state*, waiting for either s_0 or s_1 to become nonzero

Other trajectories luck out and skirt the deadlock region

Unfortunate fact: deadlock is often nondeterministic (race)

Avoiding Deadlock

Acquire shared resources in same order

```
int main()
{
    pthread_t tid[2];
    Sem_init(&mutex[0], 0, 1);  /* mutex[0] = 1 */
    Sem_init(&mutex[1], 0, 1);  /* mutex[1] = 1 */
    Pthread_create(&tid[0], NULL, count, (void*) 0);
    Pthread_create(&tid[1], NULL, count, (void*) 1);
    Pthread_join(tid[0], NULL);
    Pthread_join(tid[1], NULL);
    printf("cnt=%d\n", cnt);
    exit(0);
}
```


```
void *count(void *vargp)
{
    int i;
    int id = (int) vargp;
    for (i = 0; i < NITERS; i++) {
        P(&mutex[0]); P(&mutex[1]);
        cnt++;
        V(&mutex[id]); V(&mutex[1-id]);
    }
    return NULL;
}
```

Tid[0]:
P(s0);
P(s1);
cnt++;
V(s0);
V(s1);

Tid[1]:
P(s0);
P(s1);
cnt++;
V(s1);
V(s0);

Avoided Deadlock in Progress Graph

Thread 1



No way for trajectory to get stuck

Processes acquire locks in same order

Order in which locks released immaterial

$s_0 = s_1 = 1$

Deadlock Bugs

- Common solutions for deadlocks
 - Prevent deadlock
 - Order lock acquisition
 - Acquire lock atomically
 - Release lock voluntarily
 - Build lock-free atomic operations
 - Avoid deadlock
 - Schedule threads with global knowledge of lock acquisitions