



CS3281 / CS5281

Real-Time Scheduling

Will Hedgecock
Sandeep Neema
Bryan Ward

**Some lecture slides borrowed and adapted from the
UNC Real-Time Systems Group*



Tel (615) 343-7472 | Fax (615) 343-7440
1025 16th Avenue South Nashville, TN 37212
www.isis.vanderbilt.edu



Motivation: Cyber-Physical Systems



Surgical Robotics



Industrial Internet of Things (IIoT)



Power and Utilities



Satellites



Autonomous Vehicles

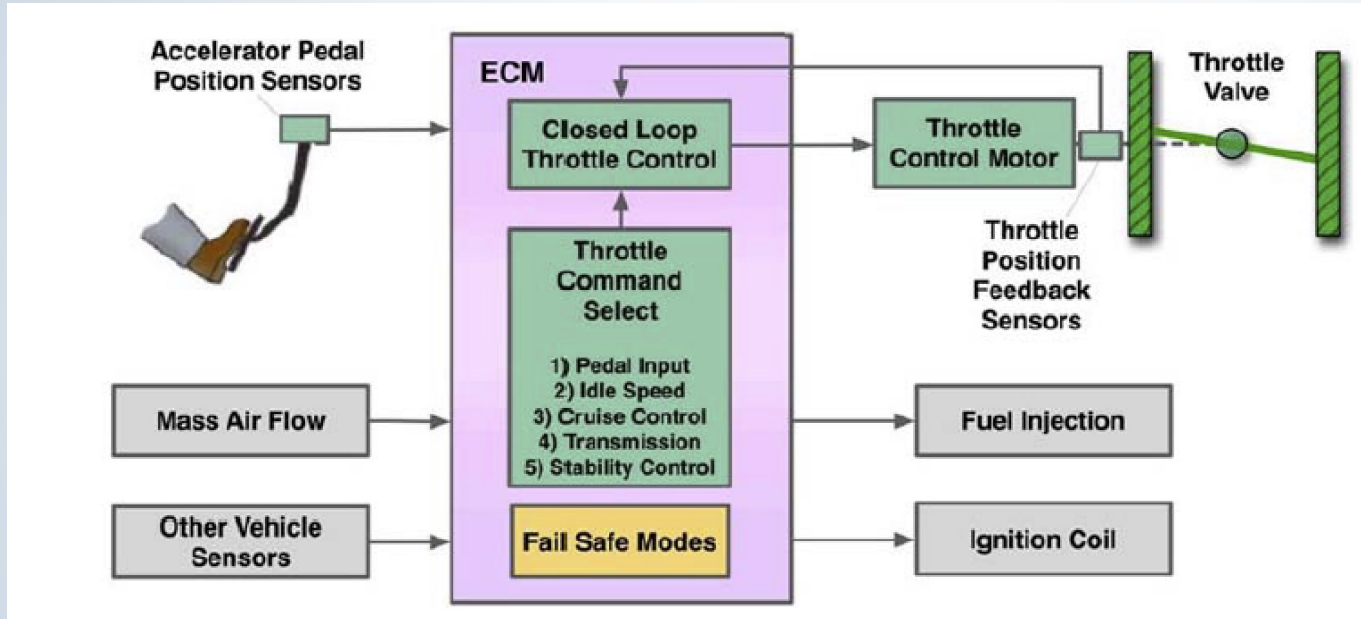


Drones & DoD Systems

Cyber Physical Systems - key characteristics

- Computing is deeply embedded and integral to safe operation of the physical system
- A large class of CPS are *safety-critical systems*
 - def: failure or malfunction may result in injury, death or severe damage to equipment/environment
- In CPS, *software-managed interactions* need to conform to laws of physics
 - that includes timing

CPS software – an example



- Electronic Throttle Control System
- Controls air + fuel + spark --> engine power
- At the heart of *Toyota's Unintended Acceleration* disaster – 89 deaths & 57 injuries

https://users.ece.cmu.edu/~koopman/pubs/koopman14_toyota_ua_slides.pdf

Real-Time Systems

Enterprise Systems




Servers, desktops, web browsing, emails, etc.

"Real Fast" Systems



Interactive processing, i.e., video games

Soft Real-Time System



Pedestrian Detection

Hard Real-Time System



Arc-Flash Relays: ~2ms to break circuit




**Interaction with the physical world requires keeping time with the real world.
Many CPS, especially safety- and mission-critical systems have strict timing requirements.**

What is a Real-Time System?

- A system with a dual notion of correctness:
 - *Logical correctness* (“it does the right thing”);
 - *Temporal correctness* (“it does it on time”).
- A system wherein *predictability* is as important as *performance*
- Real-time systems are designed based on worst case, rather than average case


Periodic Task Systems

- Set τ of periodic tasks scheduled on M cores:




Periodic Task Systems

- Set τ of periodic tasks scheduled on M cores:
 - Task $T_i = (e_i, p_i)$ releases a *job* with exec. cost e_i every p_i time units.
 - T_i 's *utilization* (or *weight*) is $u_i = e_i/p_i$.
 - *Total utilization* is $U(\tau) = \sum_{T_i} e_i/p_i$.




Periodic Task Systems

- Set τ of periodic tasks scheduled on M cores:
 - Task $T_i = (e_i, p_i)$ releases a job with exec. cost e_i every p_i time units.
 - T_i 's utilization (or weight) is $u_i = e_i/p_i$.
 - Total utilization is $U(\tau) = \sum_{T_i} e_i/p_i$.




Periodic Task Systems

- Set τ of periodic tasks scheduled on M cores:
 - Task $T_i = (e_i, p_i)$ releases a *job* with exec. cost e_i every p_i time units.
 - T_i 's *utilization* (or *weight*) is $u_i = e_i/p_i$.
 - Total utilization* is $U(\tau) = \sum_T e_i/p_i$.




Periodic Task Systems

- Set τ of periodic tasks scheduled on M cores:
 - Task $T_i = (e_i, p_i)$ releases a *job* with exec. cost e_i every p_i time units.
 - T_i 's *utilization* (or *weight*) is $u_i = e_i/p_i$.
 - *Total utilization* is $U(\tau) = \sum_{T_i} e_i/p_i$.
 - Each job of T_i has a *deadline* at the next job release of T_i .




Periodic Task Systems

- Set τ of periodic tasks scheduled on M cores:
 - Task $T_i = (e_i, p_i)$ releases a *job* with exec. cost e_i every p_i time units.
 - T_i 's *utilization* (or *weight*) is $u_i = e_i/p_i$.
 - Total utilization* is $U(\tau) = \sum_{T_i \in \tau} e_i/p_i$.
 - Each job of T_i has a *deadline* at the next job release of T_i .



Periodic Task Systems

- Set τ of periodic tasks scheduled on M cores:
 - Task $T_i = (e_i, p_i)$ releases a *job* with exec. cost e_i every p_i time units.
 - T_i 's *utilization* (or *weight*) is $u_i = e_i/p_i$.
 - Total utilization* is $U(\tau) = \sum_{T_i} e_i/p_i$.
- This is an example of an **earliest-deadline-first (EDF)** schedule.




Other Kinds of Tasks


- **Sporadic:** p_i is a minimum separation between job releases of T_i .
- For periodic or sporadic, relative deadlines d_i can be
 - **implicit:** $d_i = p_i$ (assumed unless stated otherwise);
 - **constrained:** $d_i \leq p_i$;
 - **arbitrary.**
- Also can have **aperiodic** (one-shot) jobs.
 - **hard aperiodic:** job has a deadline

Periodic vs. Sporadic

An implicit-deadline **periodic** task T_i with $p_i = 5$ and $e_i = 2$ could execute like this:



If **sporadic**, could execute like this:




Periodic vs. Sporadic


An implicit-deadline
like this:

A task system is called **synchronous** if all tasks start at time 0, and **asynchronous** otherwise.

could execute

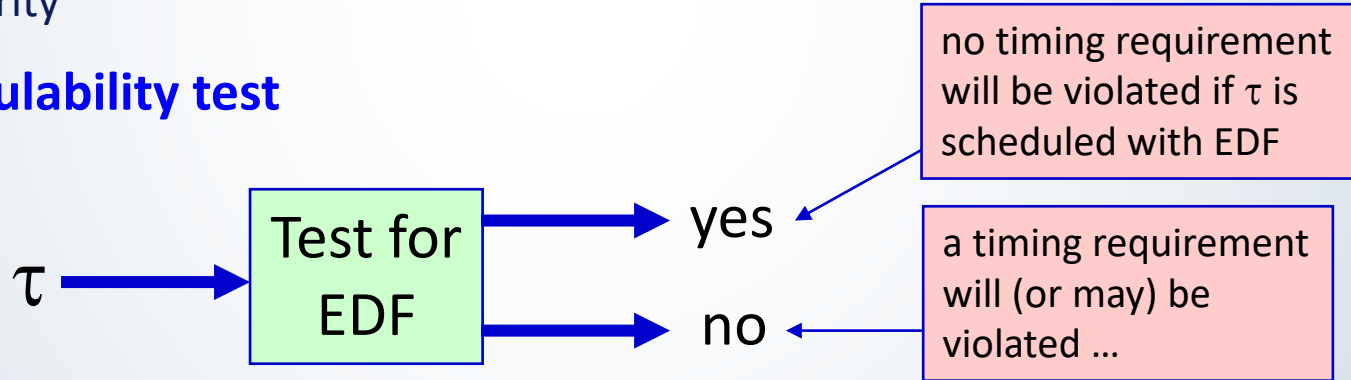


If **sporadic**, could execute like this:



Scheduling vs. Schedulability

- W.r.t. scheduling, we actually care about two kinds of algorithms:
 - **Scheduling algorithm** (of course)
 - **Example:** Earliest-deadline-first (**EDF**): Jobs with earlier deadlines have higher priority
 - **Schedulability test**



Scheduling vs. Schedulability

- W.r.t. algo
 - **Utilization loss** occurs when a test requires utilizations to be restricted to get a “yes” answer.
 - **Scheduling algorithm** (course).
 - **Example:** Earliest-deadline-first (**EDF**): Jobs with earlier deadlines have higher priority
 - **Schedulability test**
-
- ```
graph LR; T[τ] --> Test[Test for EDF]; Test --> Yes[yes]; Test --> No[no]; Yes --> BoxYes[no timing requirement will be violated if τ is scheduled with EDF]; No --> BoxNo[a timing requirement will (or may) be violated ...];
```
- The diagram illustrates the process of testing for EDF schedulability. A task  $\tau$  is input to a "Test for EDF" block. This block outputs either "yes" or "no". If the output is "yes", it indicates that no timing requirement will be violated if  $\tau$  is scheduled with EDF. If the output is "no", it indicates that a timing requirement will (or may) be violated.

# Optimality and Feasibility

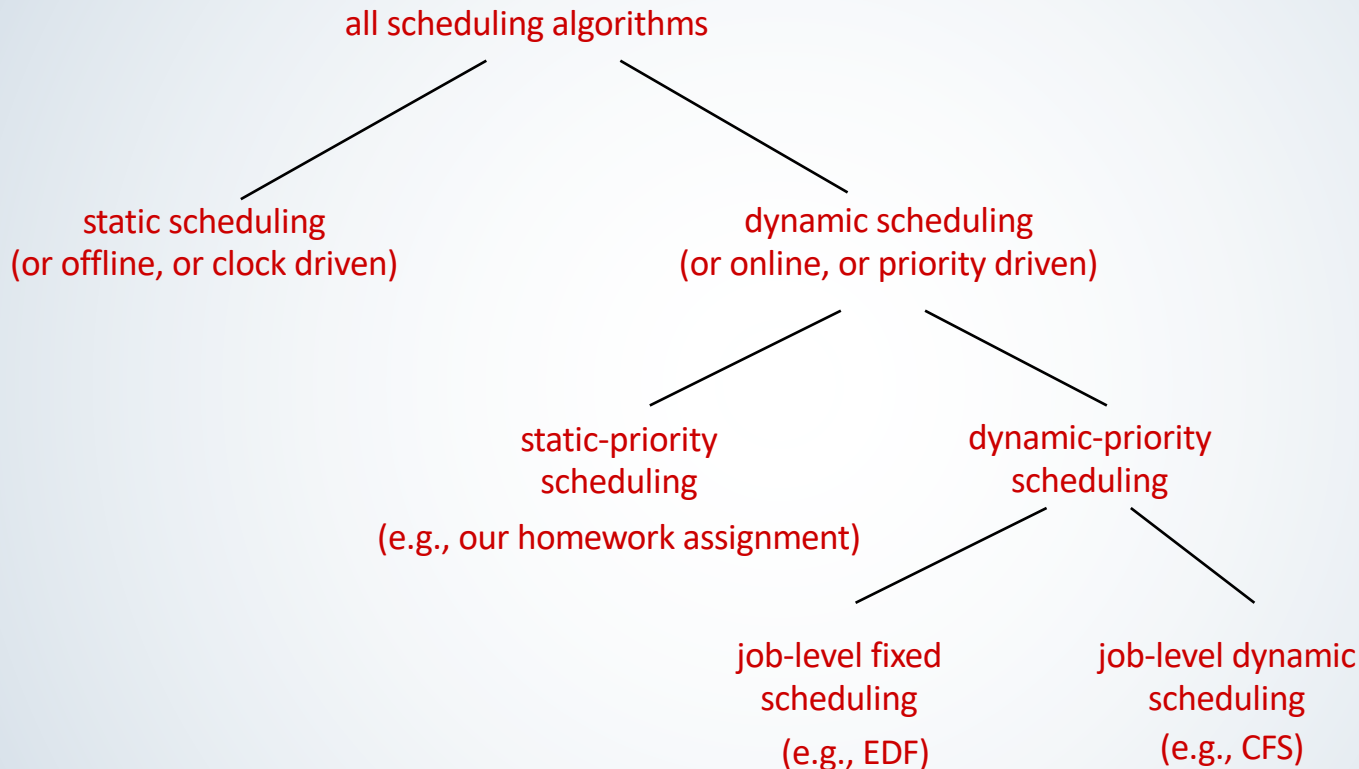
- A schedule is **feasible** if all timing constraints are met
- A task set  $T$  is **schedulable** using scheduling algorithm  $A$  if  $A$  always produces a feasible schedule for  $T$
- A scheduling algorithm is **optimal** if it always produces a feasible schedule when one exists (under any scheduling algorithm)

# Feasibility vs. Schedulability

- To most people in the real-time systems community:
  - a **feasibility test** indicates whether some algorithm from a class of algorithms can correctly schedule a given task set;
  - a **schedulability test** indicates whether a specific algorithm (e.g., EDF) can correctly schedule a given task set.
- Such tests can either be **exact** or only **sufficient**



# Classification of Scheduling Algorithms



# Static-Priority Scheduling


- Under **fixed-priority** scheduling, different jobs of a task are assigned the same priority.
- We will assume that tasks are indexed in decreasing priority order, i.e.,  $T_i$  has higher priority than  $T_k$  if  $i < k$ .
- The ready task with the highest priority is always scheduled.

# In-class examples

- Consider the following example implicit-deadline task system
  - $T_1 = (3,10)$ ,  $T_2 = (1,3)$ ,  $T_3 = (1,5)$
- Assume the task system is synchronous (i.e., all start at  $t=0$ )
- At what time does the first job of  $T_2$  complete under static-priority scheduling where  $T_1$  is the highest priority and  $T_3$  is the lowest?
- At what times does  $T_1$  complete under EDF?
  - Draw out the schedule to see

# Rate-Monotonic Scheduling

- **Priority Definition:** Tasks with smaller periods have higher priority.
- **Example Schedule:** Three tasks,  $T_1 = (0.5, 3)$ ,  $T_2 = (1, 4)$ ,  $T_3 = (2, 6)$ .



# Optimality (or not) of RM

**Theorem:** RM is not optimal.

**Exception:** RM is optimal if task set is harmonic (“simply periodic”).

## Proof:

Consider  $T_1 = (1, 2)$  and  $T_2 = (2.5, 5)$ .

Total utilization is one, so the system is feasible.

However, under RM, a deadline will be missed, regardless of how we choose to (statically) prioritize  $T_1$  and  $T_2$ .

You can work through both cases to convince yourself of this.

This proof actually shows that *no* static-priority algorithm is optimal.

# Utilization-based RM Schedulability Test

**Theorem:** [Liu and Layland] A system of  $n$  independent, preemptable sporadic tasks with relative deadlines equal to their respective periods can be feasibly scheduled on a processor according to the RM algorithm if its total utilization  $U$  is at most

$$U_{RM}(n) = n(2^{1/n} - 1)$$

Note that this is only a sufficient schedulability test.

Note: Utilization Loss =  $1 - U_{RM}(n)$



# Static-Priority Schedulability Test

(Uniprocessor, Static-Priority)

**Definition:** The time-demand function of the task  $T_i$ , denoted  $w_i(t)$ , is defined as follows.

$$w_i(t) = e_i + \sum_{k=1}^{i-1} \left\lceil \frac{t - p_k}{p_k} \right\rceil \cdot e_k \quad \text{for } 0 \leq t \leq p_i$$

**Note:** We assume tasks are indexed by priority.

For any fixed-priority **algorithm A** with  $d_i \leq p_i$  for all  $i \dots$



**Theorem:** A system  $\tau$  of sporadic, independent, preemptable tasks is schedulable on one processor by algorithm A if

$$(\forall i :: (\exists t: 0 < t \leq d_i :: w_i(t) \leq t))$$

holds.

# Real-Time Synchronization: Priority Inversions

- So far, we've assumed all jobs are independent
- A *priority inversion* occurs when a high-priority job is blocked by a low-priority one
- This is bad because HP jobs usually have more stringent timing constraints





Mars Pathfinder famously had a priority inversion when deployed and it almost caused a mission failure. A patch was sent remotely to fix it. It changed one bit.

<https://www.rapitasystems.com/blog/what-really-happened-software-mars-pathfinder-spacecraft>

# Real-Time Synchronization: Priority Inheritance

**A Common Solution:** Use *priority inheritance* (blocking job executes at blocked job's priority)

Doesn't prevent inversions but limits their duration



 Critical Section

 Priority Inversion

 Computation Outside of CS's

# Summary

- Real-time systems differ from general-purpose ones in that there exist timing requirements
- Common in cyber-physical and safety-critical systems, such as avionics, automotive, and other embedded devices
- Timing requirements inform how scheduling should be handled
- Many classes of real-time scheduling algorithms exist
- Analysis complements the scheduling implementation to prove temporal correctness