

# Lecture 13: Real-Time Scheduling

CS 3281

# Motivation – Cyber-Physical Systems



**Surgical Robotics**



**Industrial Internet of Things (IIoT)**



**Power and Utilities**



**Satellites**



**Autonomous Vehicles**



**Drones & DoD Systems**

# Real-Time Systems

Enterprise Systems




Servers, desktops, web browsing, emails, etc.

“Real Fast” Systems



Interactive processing, i.e., video games

Soft Real-Time System



Pedestrian Detection

Hard Real-Time System



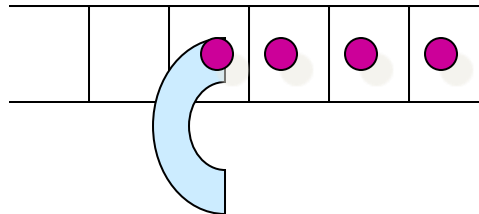
Arc-Flash Relays: ~2ms to break circuit



Interaction with the physical world requires keeping time with the real world. Many CPS, especially safety- and mission-critical systems have strict timing requirements.


# What is a Real-Time System?

- A system with a dual notion of correctness:
  - *Logical correctness* (“it does the right thing”);
  - *Temporal correctness* (“it does it on time”).
- A system wherein *predictability* is as important as *performance*.
- Real-time systems are designed based on worst case, rather than average case
- **A simple example:** A robot arm picking up objects from a conveyor belt.




# Periodic Task Systems

- Set  $\tau$  of periodic tasks scheduled on  $M$  cores:



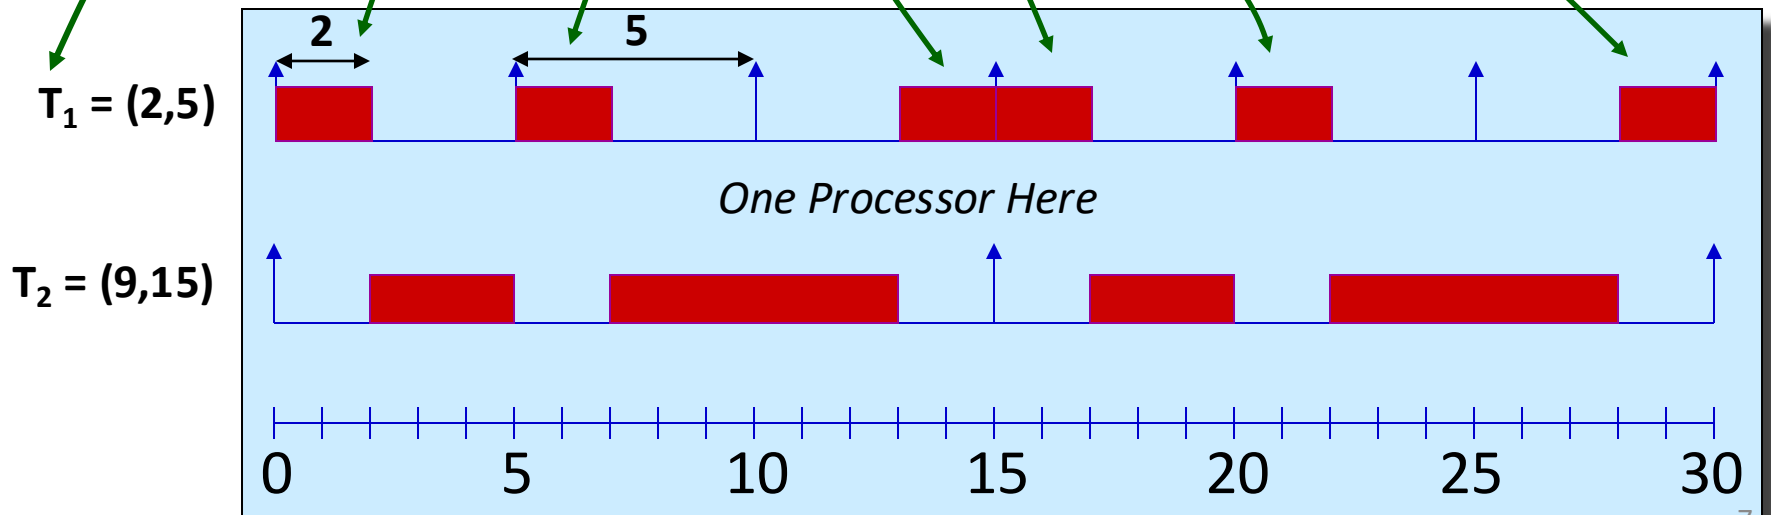
# Periodic Task Systems

- Set  $\tau$  of **periodic tasks** scheduled on  **$M$  cores**:
  - Task  $T_i = (e_i, p_i)$  releases a **job** with exec. cost  $e_i$  every  $p_i$  time units.
    - $T_i$ 's **utilization** (or **weight**) is  $u_i = e_i/p_i$ .
    - Total utilization** is  $U(\tau) = \sum_{T_i} e_i/p_i$ .




# Periodic Task Systems

- Set  $\tau$  of **periodic tasks** scheduled on  **$M$  cores**:
  - Task  $T_i = (e_i, p_i)$  releases a **job** with exec. cost  $e_i$  every  $p_i$  time units.
    - $T_i$ 's **utilization** (or **weight**) is  $u_i = e_i/p_i$ .
    - Total utilization** is  $U(\tau) = \sum_T e_i/p_i$ .



# Periodic Task Systems


- Set  $\tau$  of **periodic tasks** scheduled on  **$M$  cores**:
  - Task  $T_i = (e_i, p_i)$  releases a **job** with exec. cost  $e_i$  every  $p_i$  time units.
    - $T_i$ 's **utilization** (or **weight**) is  $u_i = e_i/p_i$ .
    - Total utilization** is  $U(\tau) = \sum_{T_i \in \tau} e_i/p_i$ .






# Periodic Task Systems

- Set  $\tau$  of **periodic tasks** scheduled on **M cores**:
  - Task  $T_i = (e_i, p_i)$  releases a **job** with exec. cost  $e_i$  every  $p_i$  time units.
    - $T_i$ 's **utilization** (or **weight**) is  $u_i = e_i/p_i$ .
    - Total utilization** is  $U(\tau) = \sum_{T_i} e_i/p_i$ .
  - Each job of  $T_i$  has a **deadline** at the next job release of  $T_i$ .



# Periodic Task Systems


- Set  $\tau$  of **periodic tasks** scheduled on  **$M$  cores**:
  - Task  $T_i = (e_i, p_i)$  releases a **job** with exec. cost  $e_i$  every  $p_i$  time units.
    - $T_i$ 's **utilization** (or **weight**) is  $u_i = e_i/p_i$ .
    - Total utilization** is  $U(\tau) = \sum_{T_i} e_i/p_i$ .
  - Each job of  $T_i$  has a **deadline** at the next job release of  $T_i$ .



# Periodic Task Systems

- Set  $\tau$  of **periodic tasks** scheduled on **M cores**:
  - Task  $T_i = (e_i, p_i)$  releases a **job** with exec. cost  $e_i$  every  $p_i$  time units.
    - $T_i$ 's **utilization** (or **weight**) is  $u_i = e_i/p_i$ .
    - Total utilization** is  $U(\tau) = \sum_{T_i} e_i/p_i$ .

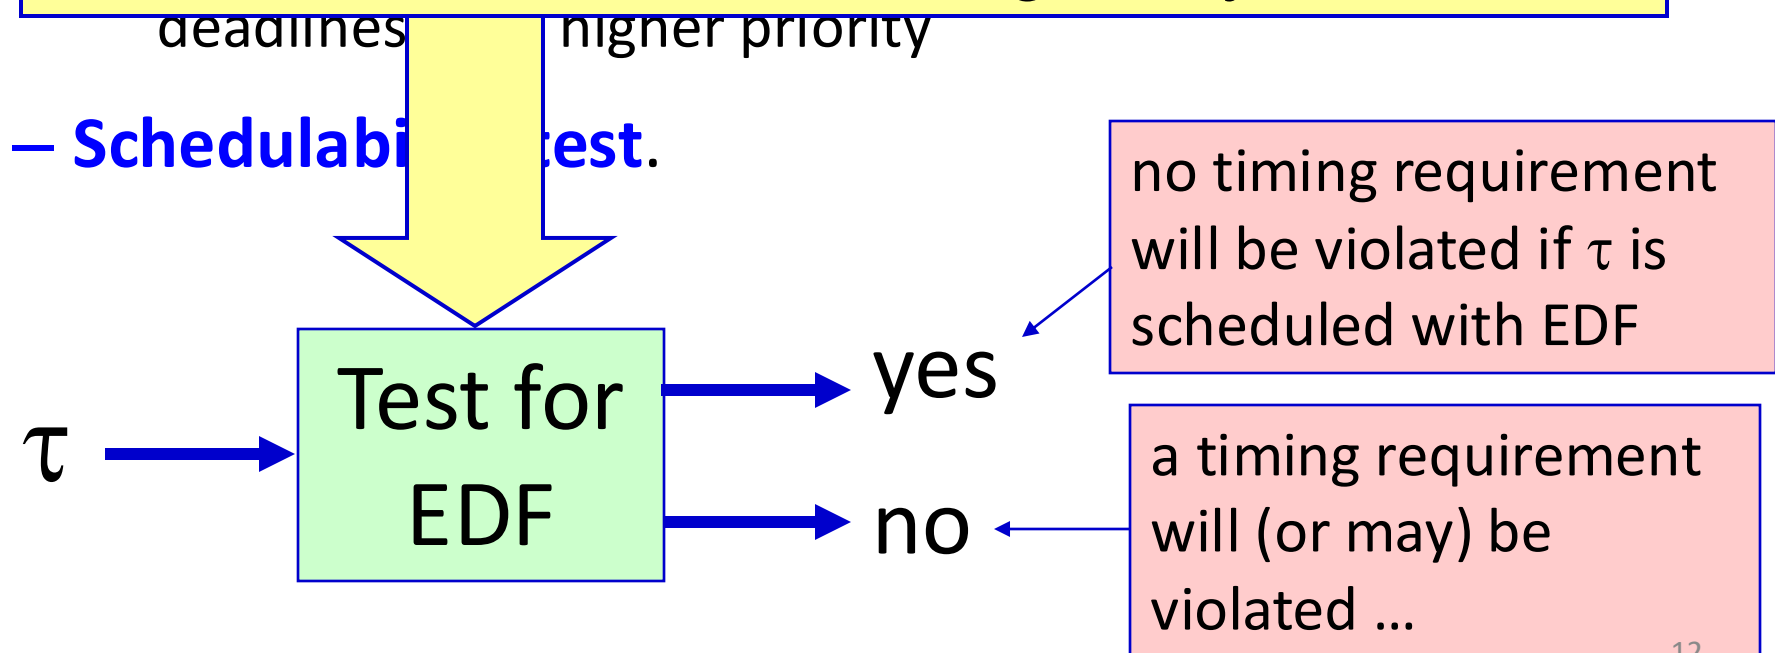
This is an example of an **earliest-deadline-first (EDF)** schedule.



# Scheduling vs. Schedulability

- W.r.t. scheduling, we actually care about two kinds of algorithms:

**Utilization loss** occurs when a test requires utilizations to be restricted to get a “yes” answer.



# Optimality and Feasibility

- A schedule is **feasible** if all timing constraints are met
- A task set  $\tau$  is **schedulable** using scheduling algorithm A if A produces a feasible schedule for  $\tau$
- A scheduling algorithm is **optimal** if it provides a feasible schedule for a schedulable task set

# Static-Priority Scheduling


- Under **fixed-priority** scheduling, different jobs of a task are assigned the same priority.
- We will assume that tasks are indexed in decreasing priority order, i.e.,  **$T_i$  has higher priority than  $T_k$  if  $i < k$ .**
- The ready task with the highest priority is always scheduled.

# Rate-Monotonic Scheduling

(Liu and Layland)

**Priority Definition:** Tasks with smaller periods have higher priority.

**Example Schedule:** Three tasks,  $T_1 = (0.5, 3)$ ,  $T_2 = (1, 4)$ ,  $T_3 = (2, 6)$ .



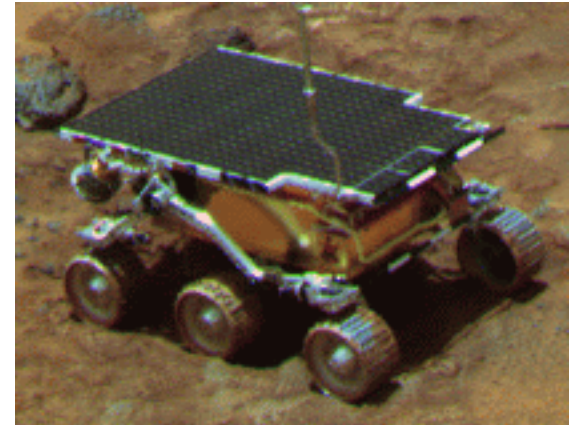
# RT Synchronization 101

## Priority Inversions

So far we've assumed all jobs are independent.

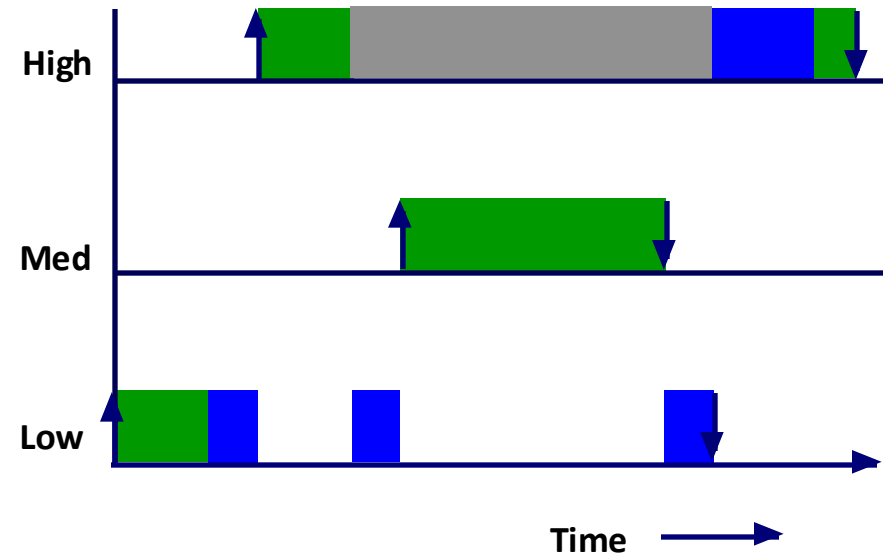
A *priority inversion* occurs when a high-priority job is blocked by a low-priority one.

This is bad because HP jobs usually have more stringent timing constraints.



Mars Pathfinder infamously had a priority inversion when deployed and it almost caused a mission failure. A patch was sent remotely patched to fix.

<https://www.rapitasystems.com/blog/what-really-happened-software-mars-pathfinder-spacecraft>



Critical Section



Priority Inversion



Computation Outside of CS's

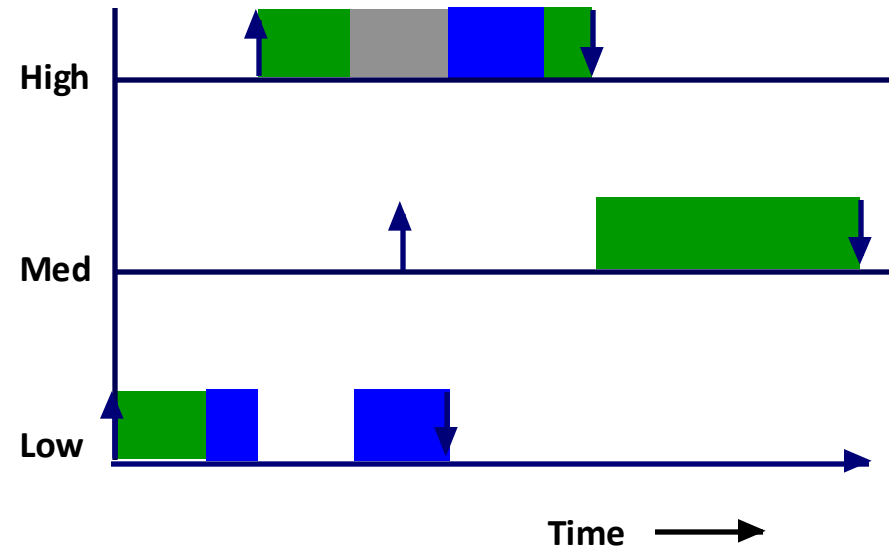
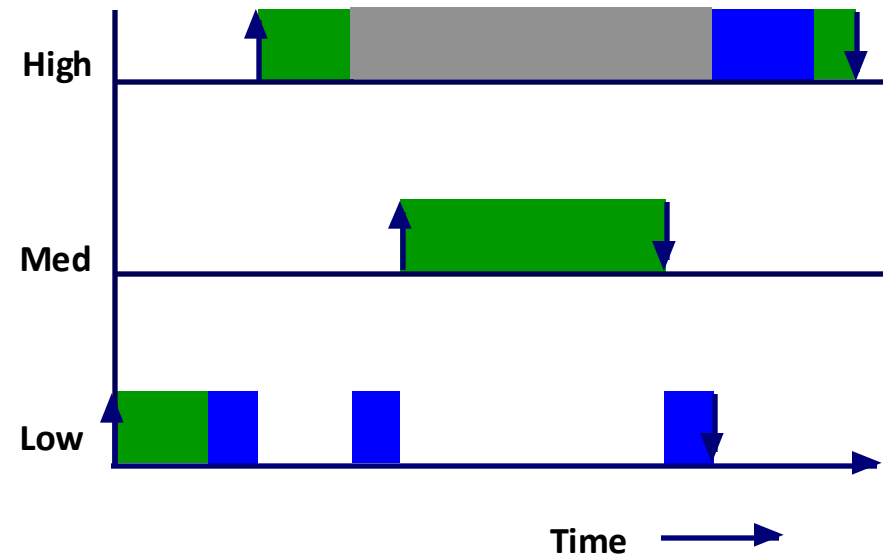
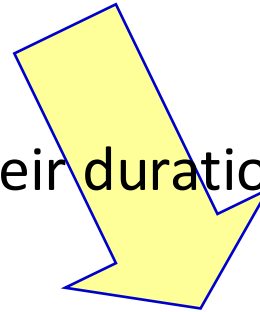


# RT Synchronization 101

## Priority Inheritance

**A Common Solution:** Use *priority inheritance* (blocking job executes at blocked job's priority).

Doesn't prevent inversions but limits their duration.



Critical Section



Priority Inversion



Computation Outside of CS's

# Scheduler Classes

- Linux has different algorithms for scheduling different types of processes
  - Called scheduler classes
- Each class implements a different but “pluggable” algorithm for scheduling
  - Within a class, you can set the policy
- RT class: `SCHED_FIFO`, `SCHED_RR`
- Non-RT class: `SCHED_OTHER` (CFS)

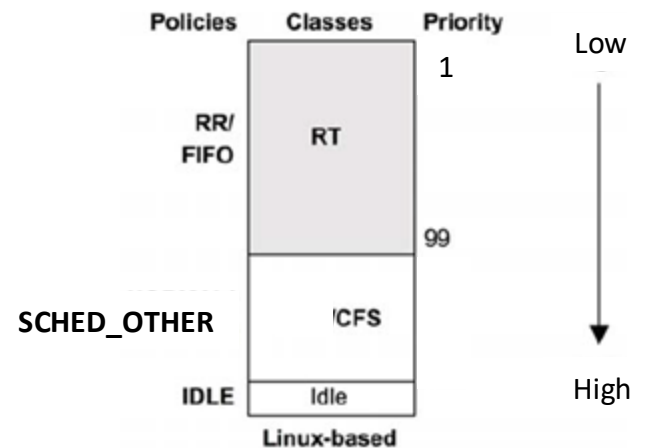


Figure from, “*Systems Performance: Enterprise and the Cloud*” by Brendan Gregg

# Exercise

Consider two processes, T1 and T2, where  $p_1 = 50$ ,  $e_1 = 25$ ,  $p_2 = 75$ ,  $e_2 = 30$ .

Illustrate the scheduling of these two processes using

- earliest-deadline-first (EDF) scheduling
- rate-monotonic scheduling

# Exercise

Let A, B, and C be three tasks. A has the highest priority and C has the lowest priority. A and C use a shared resource protected by a mutex.

Suppose a scheduler makes decisions about scheduling tasks based on their priorities. That is, the scheduler runs, among tasks, the one with the highest priority. Such a task runs to completion or blocks for any reason.

- What are the completion times of A, B, C before applying the priority inheritance?
- What are the completion times of A, B, C after applying the priority inheritance?

Task	Start time	Before lock	Critical section	After lock
A	2	1	2	1
B	3.8	4.2		
C	0	1	3	1

# Summary

- Real-time systems differ from general-purpose ones in that there exist timing requirements
- Common in cyber-physical and safety-critical systems, such as avionics, automotive, and other embedded devices.
- Timing requirements inform how scheduling should be handled
- Many classes of real-time scheduling algorithms
- Analysis complements the scheduling implementation to prove temporal correctness