



CS3281 / CS5281  
**I/O Devices**

Will Hedgecock  
Sandeep Neema  
Bryan Ward

*\*Some lecture slides borrowed and adapted from  
Andrea Arpaci-Dusseau*



# Motivation

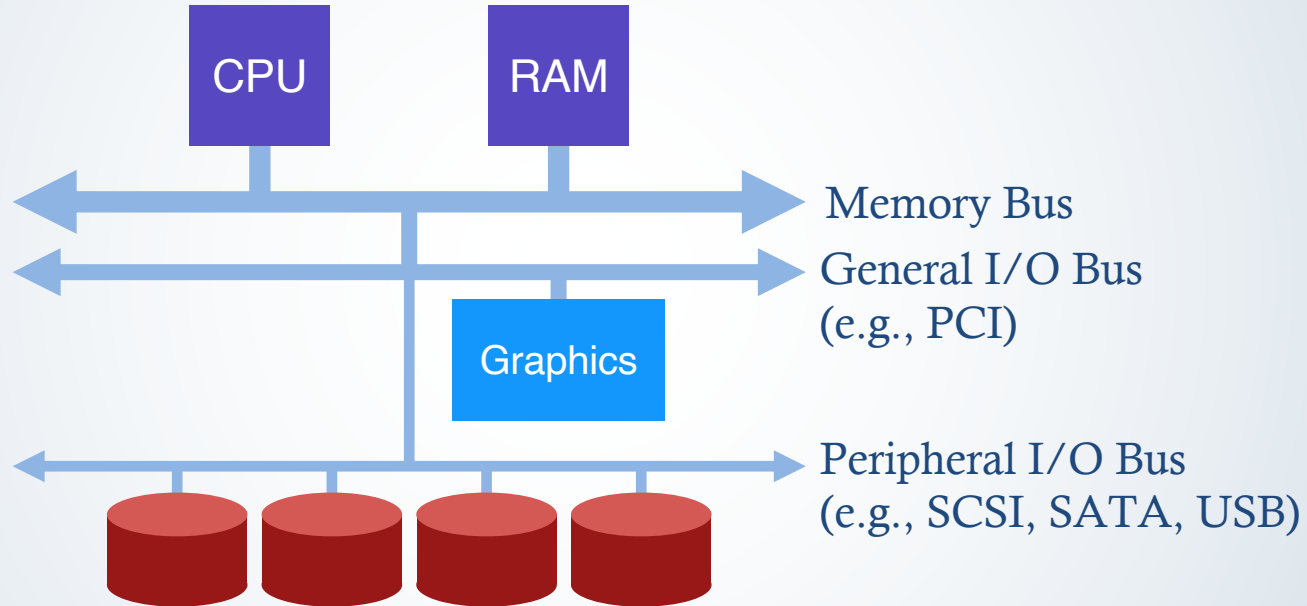
What good is a computer without any I/O devices?

- e.g., keyboard, display, disks

We want:

- **H/W** that will let us plug in different devices
- **OS** that can interact with different combinations
- I/O also allows for *persistence*
  - RAM is *volatile*, i.e., contents are lost when the machine restarts

# Hardware support for I/O



Why use hierarchical buses?

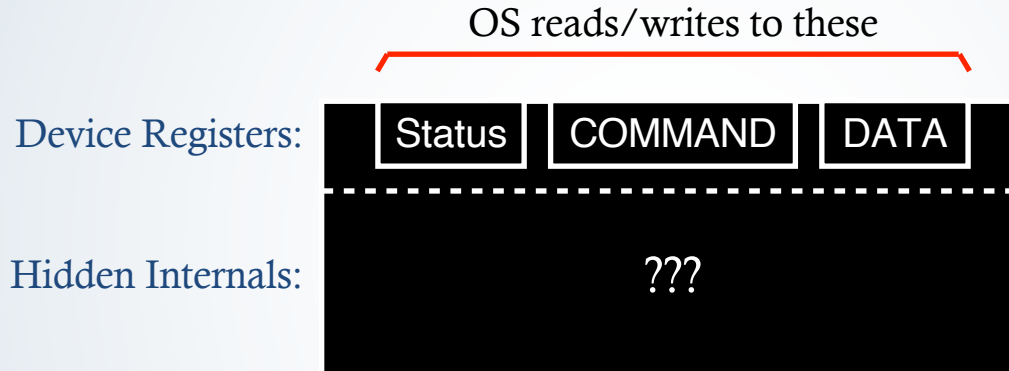
# Canonical Device

OS reads/writes to these

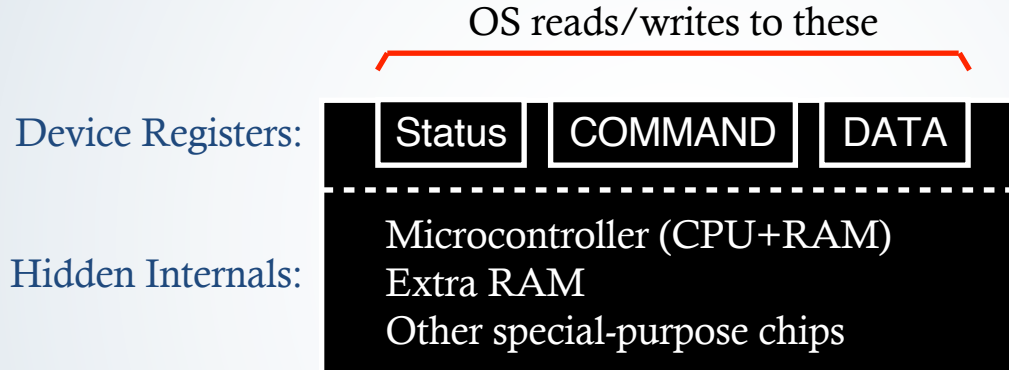
Device Registers:




# Canonical Device



# Canonical Device



# Example Write Protocol



```
while (STATUS == BUSY)
```

```
    ; // spin
```

```
Write data to DATA register
```

```
Write command to COMMAND register
```

```
while (STATUS == BUSY)
```

```
    ; // spin
```

This is called polling when the processor “asks” what the hardware is doing, often continuously

CPU:

Disk:

```
while (STATUS == BUSY)           // 1
    ;
Write data to DATA register      // 2
Write command to COMMAND register // 3
while (STATUS == BUSY)           // 4
    ;
```



CPU: A

Disk: C

```
while (STATUS == BUSY)           // 1
    ;
Write data to DATA register      // 2
Write command to COMMAND register // 3
while (STATUS == BUSY)           // 4
    ;
```

A wants to do I/O



CPU: A

Disk: C

```
while (STATUS == BUSY)           // 1
```


```
;
```

```
Write data to DATA register     // 2
```

```
Write command to COMMAND register // 3
```


```
while (STATUS == BUSY)           // 4
```

```
;
```




```
while (STATUS == BUSY)           // 1
;

Write data to DATA register      // 2
Write command to COMMAND register // 3
while (STATUS == BUSY)           // 4
;
```




```
while (STATUS == BUSY)           // 1
;

Write data to DATA register      // 2
Write command to COMMAND register // 3
while (STATUS == BUSY)           // 4
;
```




```
while (STATUS == BUSY)           // 1
;

Write data to DATA register      // 2
Write command to COMMAND register // 3
while (STATUS == BUSY)           // 4
;
```



```
while (STATUS == BUSY)           // 1
;

Write data to DATA register      // 2
Write command to COMMAND register // 3
while (STATUS == BUSY)           // 4
;
```




```
while (STATUS == BUSY)           // 1
;

Write data to DATA register      // 2

Write command to COMMAND register // 3

while (STATUS == BUSY)           // 4
;
```

how to avoid spinning?  
interrupts!



```
while (STATUS == BUSY)                // 1
    wait for interrupt;


Write data to DATA register           // 2

Write command to COMMAND register      // 3

while (STATUS == BUSY)                // 4
    wait for interrupt;
```



how to avoid spinning?  
interrupts!



```
while (STATUS == BUSY)                // 1
```

```
    wait for interrupt;
```

```
Write data to DATA register           // 2
```

```
Write command to COMMAND register     // 3
```

```
while (STATUS == BUSY)                // 4
```

```
    wait for interrupt;
```

# Interrupts vs. Polling

Are interrupts ever worse than polling?

Fast device: Better to spin than take interrupt and context-switching overhead

- Device time unknown? Hybrid approach (spin then use interrupts)


Flood of interrupts arrive

- Can lead to livelock (always handling interrupts)
- Better to ignore interrupts while making some progress handling them


Other improvement

- Interrupt coalescing (hardware batches together several interrupts)
- Split interrupt handling
  - Top half: respond to the interrupt and store relevant data
  - Bottom half: do the real interrupt-handling work at a later point

# Protocol Variants



- **Status checks:** polling vs. interrupts
- **Data:** programmed I/O (PIO) vs. direct memory access (DMA)
- **Control:** special instructions vs. memory-mapped I/O



```
while (STATUS == BUSY)                // 1
    wait for interrupt;

Write data to DATA register           // 2

Write command to COMMAND register      // 3

while (STATUS == BUSY)                // 4
    wait for interrupt;
```

what else can we optimize?

data transfer!


# Programmed I/O vs. Direct Memory Access

## **PIO** (Programmed I/O):

- CPU directly tells device what the data is
- One instruction for each byte/word
- Efficient for a few bytes/words, but *scales terribly*

## **DMA** (Direct Memory Access):

- CPU leaves data in memory
- Device reads/writes data directly from/to memory
- One instruction to send a pointer to the data to send
- Efficient for large data transfers




```
while (STATUS == BUSY)                // 1
    wait for interrupt;


Write data to DATA register           // 2

Write command to COMMAND register      // 3

while (STATUS == BUSY)                // 4
    wait for interrupt;
```




```
while (STATUS == BUSY)                // 1
    wait for interrupt;
Write data to DATA register        // 2
Write command to COMMAND register      // 3
while (STATUS == BUSY)                // 4
    wait for interrupt;
```



```
while (STATUS == BUSY) // 1
    wait for interrupt;
Write data to DATA register // 2
Write command to COMMAND register // 3
while (STATUS == BUSY) // 4
    wait for interrupt;
```




# Protocol Variants



**Status checks:** polling vs. interrupts

**Data:** PIO vs. DMA

**Control:** special instructions vs. memory-mapped I/O



```
while (STATUS == BUSY)                // 1
    wait for interrupt;

Write data to DATA register         // 2

Write command to COMMAND register    // 3

while (STATUS == BUSY)                // 4
    wait for interrupt;
```

how does OS read and write registers?

# Special Instructions vs. Mem-Mapped I/O

## Special instructions


- each device has a port
- in/out instructions (x86) communicate with device

## Memory-Mapped I/O

- H/W maps registers into address space
- loads/stores sent to device

Doesn't matter much (both are used)

# Protocol Variants



**Status checks:** polling vs. interrupts

**Data:** PIO vs. DMA

**Control:** special instructions vs. memory-mapped I/O

# Variety is a Challenge

## Problem:

- many, many devices
- each has its own protocol

How can we avoid writing a slightly different OS for each H/W combination?

Write device driver for each device

Drivers are **70%** of Linux source code

# Storage Stack

application

.....  
file system

.....  
scheduler

.....  
driver

.....  
hard drive

build common interface  
on top of all HDDs